

Categorical predictors

WiSe23/24

Daniela Palleschi

Humboldt-Universität zu Berlin

2023-10-11

Learning Objectives

Today we will learn...

- about categorical predictors
- how to interpret different contrast coding

Set-up environment

```
1 # suppress scientific notation
2 options(scipen=999)
```

```
1 # load libraries
2 pacman::p_load(
3     tidyverse,
4     here,
5     broom,
6     lme4,
7     janitor,
8     languageR)
```

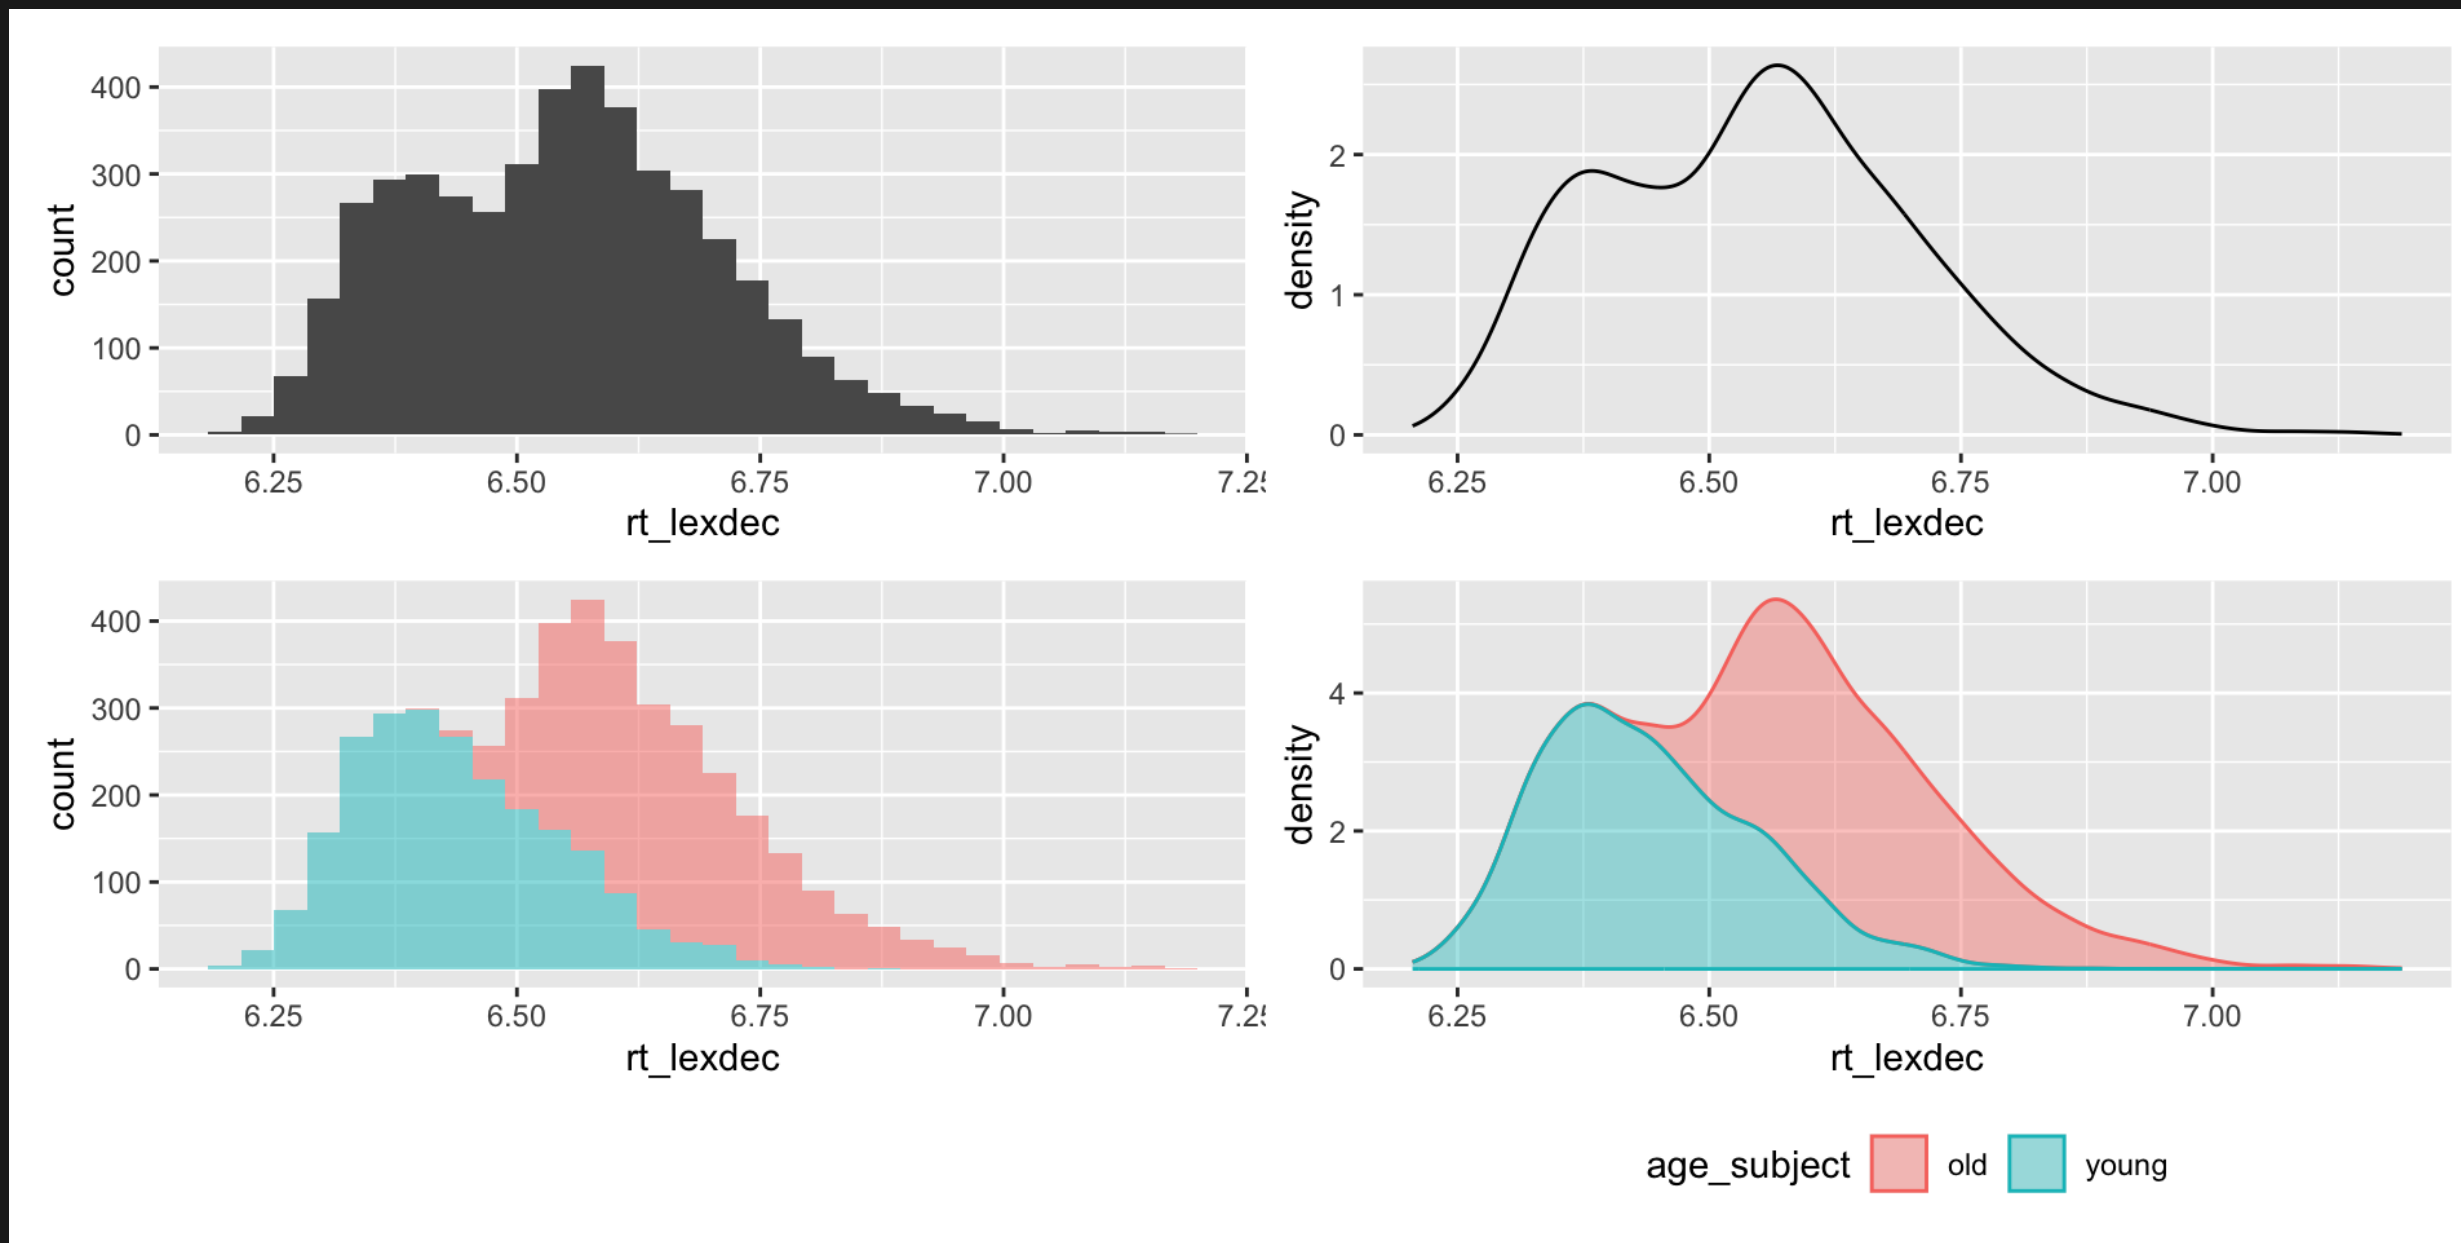
Load data

- load in the the dataset from the **languageR** package

```
1 df_freq_eng <-  
2   as.data.frame(english) |>  
3   # keep relevant variables  
4   dplyr::select(RTlexdec, RTnaming, Word, LengthInLetters, AgeSubject, WrittenFrequency) |>  
5   # rename some variables  
6   rename(rt_lexdec = RTlexdec,  
7          rt_naming = RTnaming,  
8          freq_written = WrittenFrequency) |>  
9   clean_names() |>  
10  # standardize continuous predictors  
11  mutate(  
12    freq_z = scale(freq_written),  
13    length_z = scale(length_in_letters)  
14  ) |>  
15  # move 'word' to front  
16  relocate(word) |>  
17  # arrange alphabetically by 'word'  
18  arrange(word)
```

Bimodal distribution

- in your exploratory data analysis, you might've noticed a *bimodal* distribution.



Bimodal distribution

- this is a *bimodal* distribution
 - there are two *modes* (most frequent value, i.e., peak in a histogram)
- We know that there were two subject groups: old and young
 - it might be that each group has a different mode

Re-run our model

- re-run our multiple regression model (reaction times ~ frequency + length)

```
1 fit_freq_length <-  
2   lm(rt_lexdec ~ freq_z*length_z,  
3     data = df_freq_eng)
```


Model fit and overfitting

```
1 glance(fit_freq_length)$r.squared
```

```
[1] 0.1896649
```

```
1 glance(fit_freq_length)$adj.r.squared
```

```
[1] 0.1891323
```

- seems like we don't have any overfitting in our model (R^2 and adjusted R^2 are comparable)

Model coefficients

- look at our coefficients.

```
1 tidy(fit_freq_length) |> select(term, estimate)
```

```
# A tibble: 4 × 2
  term          estimate
<chr>          <dbl>
1 (Intercept)    6.55
2 freq_z        -0.0682
3 length_z       0.00328
4 freq_z:length_z -0.00196
```

- looks similar to the dataset we explored yesterday
- the bimodal distribution we saw earlier suggests age group could be an important predictor
- does the effect of frequency and length also differ as a function of age group?

Categorical predictors

- we'd predict longer reading times for older participants than younger participants
 - although we should hypothesise *before* collecting and visualising our data!
- though age is numerical, all we have is two categories: old or young

Including a categorical predictor

- include `age_subject` in our model

```
1 fit_age <-  
2   lm(rt_lexdec ~ freq_z*length_z + age_subject,  
3     data = df_freq_eng)
```

Model fit

- compare R^2 and adjusted R^2
- R^2 our model without age as a predictor:

```
1 # rt_lexdec ~ freq_z*length_z
2 glance(fit_freq_length)$adj.r.squared
```

```
[1] 0.1891323
```

- R^2 our model with age as a predictor:

```
1 # rt_lexdec ~ freq_z*length_z + age_subject
2 glance(fit_age)$r.squared
```

```
[1] 0.6888949
```

- adjusted R^2 our model with age as a predictor:

```
1 # rt_lexdec ~ freq_z*length_z + age_subject
2 glance(fit_age)$adj.r.squared
```

```
[1] 0.6886222
```

- large increase in proportion of variance explained when we include age
- and the R^2 and adjusted R^2 values are comparable for the model with age
- this suggests that age captures variance that was not explained without it

Check for absence of collinearity

```
1 car::vif(fit_age)
```

freq_z	length_z	age_subject	freq_z:length_z
1.012553	1.004461	1.000000	1.008108

- VIF values for all coefficients are near 1
 - this indicates that our predictors all contribute to the variance explained by the model and are not correlated

Contrasts

- let's take a look at our model estimates

```
1 tidy(fit_age) |> select(term, estimate)
```

```
# A tibble: 5 × 2
```

	term	estimate
	<chr>	<dbl>
1	(Intercept)	6.66
2	freq_z	-0.0682
3	length_z	0.00328
4	age_subjyoung	-0.222
5	freq_z:length_z	-0.00196

- there is a negative slope for **age_subjyoung**
 - reaction times decrease when...what?
- how does a categorical variable get fit to a line?
- the factor levels (i.e., the categories in a categorical variable) are given numerical values
 - We call these numerical values mapped onto factor levels contrast coding

Dummy coding/treatment contrasts

- we can check the contrasts with `contrasts()`

```
1 contrasts(df_freq_eng$age_subject)
```

	young
old	0
young	1

- `old` was coded at 0 and `young` as 1
- our slope for `age_subjectyoung` represents the change in reaction times when we move from `old` to `young`
- this is called `treatment coding` (a.k.a., dummy coding), where one factor level is coded as 0 and the other as 1

Age-only model

- remove frequency and length to focus on **age_subject**
- use raw reaction times, to more easily interpret the results

```
1 fit_age <-  
2   lm(exp(rt_lexdec) ~ age_subject,  
3     data = df_freq_eng)
```

- what's the variance explained by our (simple) model with only age as a predictor?

```
1 glance(fit_age)$r.squared  
[1] 0.4682224
```

- R^2 is lower than when we included frequency and length
 - but higher than our model with frequency and length but no age

Age-only coefficients

```
1 tidy(fit_age) |> select(term, estimate)
```

```
# A tibble: 2 × 2
  term          estimate
<chr>         <dbl>
1 (Intercept)    787.
2 age_subjectyoung -157.
```

- reaction times decrease by 157ms going from old to young group compared to the old group
- what does the intercept represent here?

```
1 df_freq_eng |>
2   select(rt_lexdec, age_subject) |>
3   mutate(rt_lexdec = exp(rt_lexdec)) |>
4   summary()
```

rt_lexdec	age_subject
Min. : 495.4	old :2284
1st Qu.: 617.4	young:2284
Median : 699.6	
Mean : 708.1	
3rd Qu.: 775.3	
Max. :1323.2	

- don't see the intercept value there

Summarising group effects

- our intercept was 786.72, but that wasn't the grand mean reaction time
 - what is the intercept?
- how does `rt_lexdec` look for the two groups?

```
1 df_freq_eng |>
2   select(rt_lexdec, age_subject) |>
3   mutate(rt_lexdec = exp(rt_lexdec)) |>
4   summarise(mean = mean(rt_lexdec),
5             min = min(rt_lexdec),
6             max = max(rt_lexdec),
7             .by = "age_subject"
8   )
```

	age_subject	mean	min	max
1	young	629.5473	495.38	971.8
2	old	786.7200	603.77	1323.2

- the intercept corresponds to the mean reaction time for the old group. Why?
 - because `old` coded as 0

Intercept at 0

- the intercept corresponds to the value of y when x is 0
 - when predictors are *centered*, this will correspond to the mean value of y , because when $x = 0$ it aligns with the centre value of y
 - when predictors are not centered, this will correspond to the value of y when x is 0 in the original unit of measurement

Default contrasts

- which variable is coded as 0?
 - R simply takes the first level name alphabetically: **old** comes before **young**, so **old** was automatically taken as the 'baseline' to which **young** was compared
- if we were to add the slope to the intercept, we would get the mean for the young group. Why is this?

```
1 coef(fit_age)[ '(Intercept)' ] + coef(fit_age)[ 'age_subjctyoung' ]
```

```
(Intercept)  
629.5473
```

Simple linear regression as a two-sample t-test

- this actually is the same thing as a *t*-test:

```
1 t.test(exp(rt_lexdec) ~ age_subject, data = df_freq_eng)
```

Welch Two Sample t-test

data: exp(rt_lexdec) by age_subject

t = 63.406, df = 4144.6, p-value < 0.000000000000000022

alternative hypothesis: true difference in means between group old and group young is not equal to 0

95 percent confidence interval:

152.3128 162.0325

sample estimates:

mean in group old mean in group young

786.7200

629.5473

- if we compare this to our model, we see that the *t*- and *p*-values are identical (more on these later).

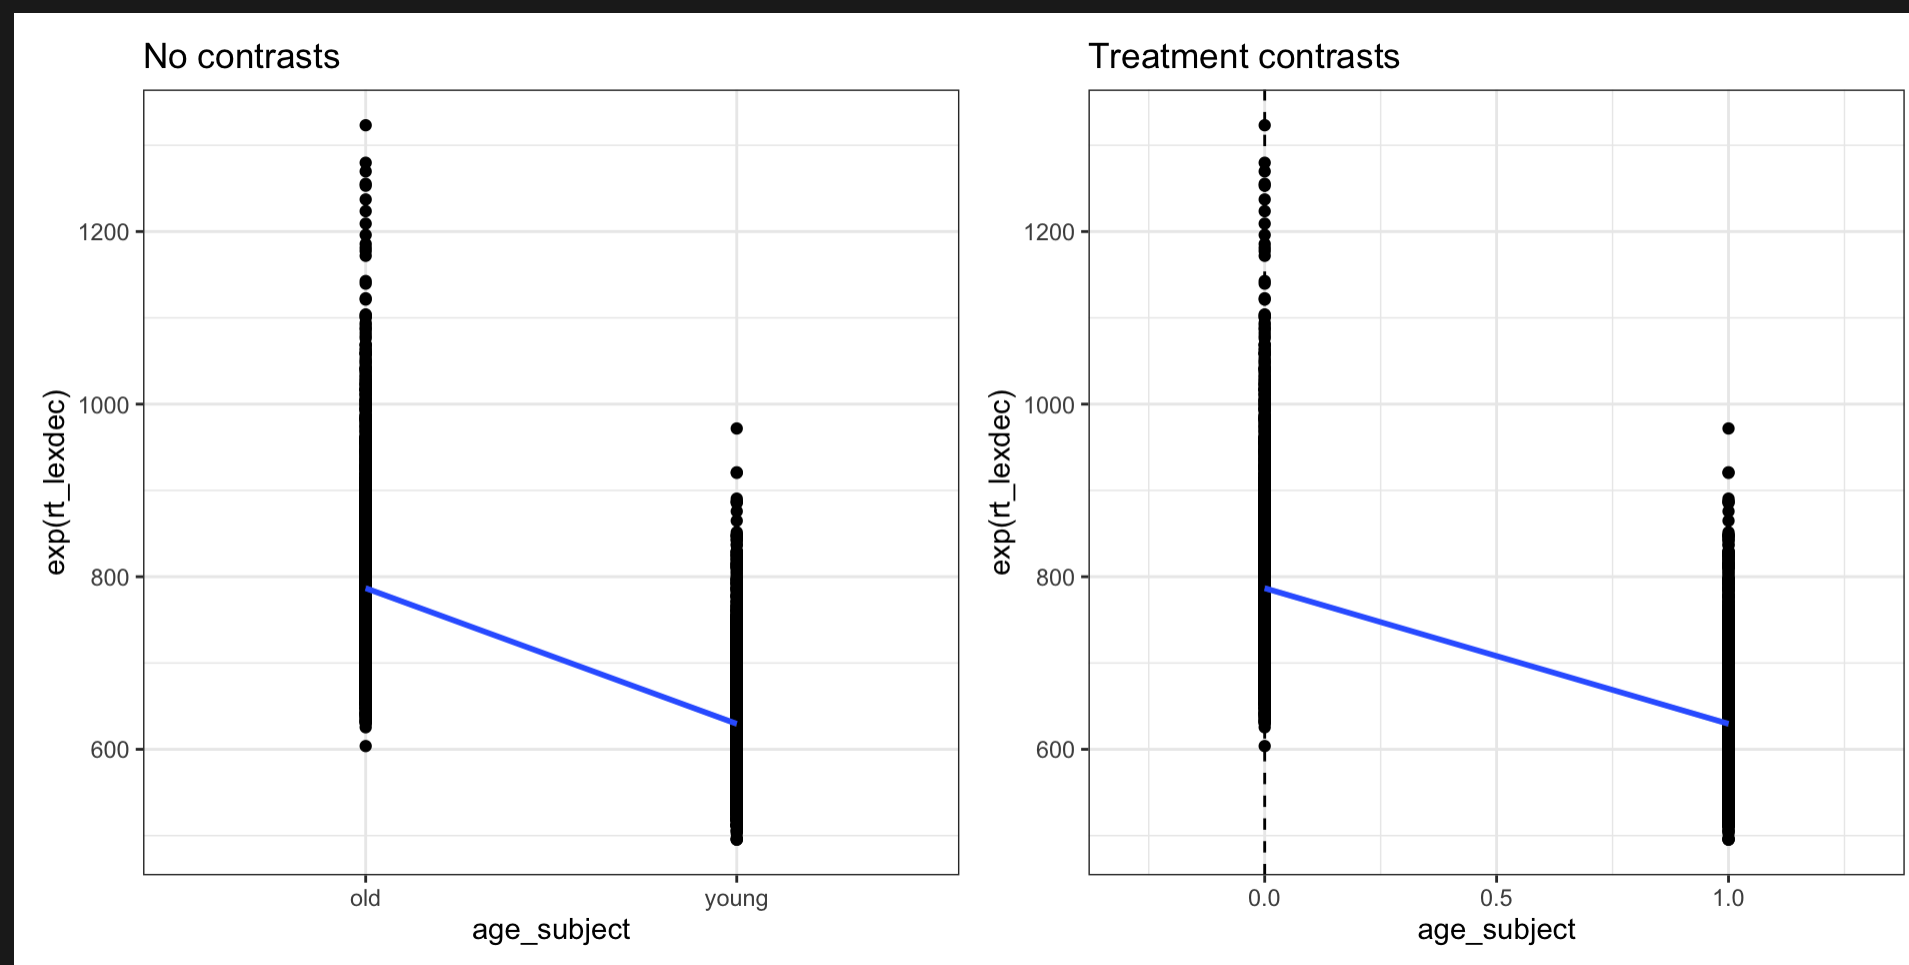
```
1 tidy(fit_age)
```

A tibble: 2 × 5

	term <chr>	estimate <dbl>	std.error <dbl>	statistic <dbl>	p.value <dbl>
1	(Intercept)	787.	1.75	449.	0
2	age_subjectyoung	-157.	2.48	-63.4	0

Visualizing treatment contrasts

► Code



Sum contrasts/coding

- sum coding is another frequently used coding scheme
 - essentially centring categorical variables
- simplifies interpretation of interaction effects
- instead of 0 and 1, we set our contrasts to +/-1 or 0.5 (I prefer 0.5)

Setting sum contrasts

- ensure we're working with a factor

```
1 # first, make sure your variable is a factor
2 df_freq_eng$age_subject <- as.factor(df_freq_eng$age_subject)
```

- check it is a factor (could do this first)

```
1 # check
2 class(df_freq_eng$age_subject)
```

```
[1] "factor"
```

contr.sum()

- we can use the `contr.sum()` function to set sum contrasts
 - takes as its argument the number of factor levels

```
1 # next, you could use the contr.sum() function
2 contrasts(df_freq_eng$age_subject) <- contr.sum(2) # where 2 means we have 2 levels
3 contrasts(df_freq_eng$age_subject)
```

```
      [,1]
old       1
young    -1
```

- `old` is coded as -1 and `young` as $+1$
- I prefer to use ± 0.5 for reasons we don't need to go into here
 - I would also prefer to have `young` coded in the negative value, and `old` in the positive value
 - this aids in the way I interpret the slope: a change in reaction times for the older group compared to the younger group

By-hand

```
1 #or, you could manually control the sum contrasts
2 # check the order of the levels
3 levels(df_freq_eng$age_subject)
```

```
[1] "old"    "young"
```

```
1 # code 'old' as +.5 and 'young' as -.5
2 contrasts(df_freq_eng$age_subject) <- c(+0.5, -0.5)
3 contrasts(df_freq_eng$age_subject)
```

```
      [,1]
old      0.5
young -0.5
```

Model with sum coded factor

- run our model

```
1 fit_age_sum <-  
2   lm(exp(rt_lexdec) ~ age_subject,  
3     data = df_freq_eng)
```

```
1 glance(fit_age_sum)$r.squared
```

```
[1] 0.4682224
```

```
1 glance(fit_age)$r.squared
```

```
[1] 0.4682224
```

- no difference in variance account for by our model (remember, centering a variable just shifts values, doesn't affect the relationship between values)

Coefficients

```
1 tidy(fit_age_sum) |> select(term, estimate)
```

```
# A tibble: 2 × 2
```

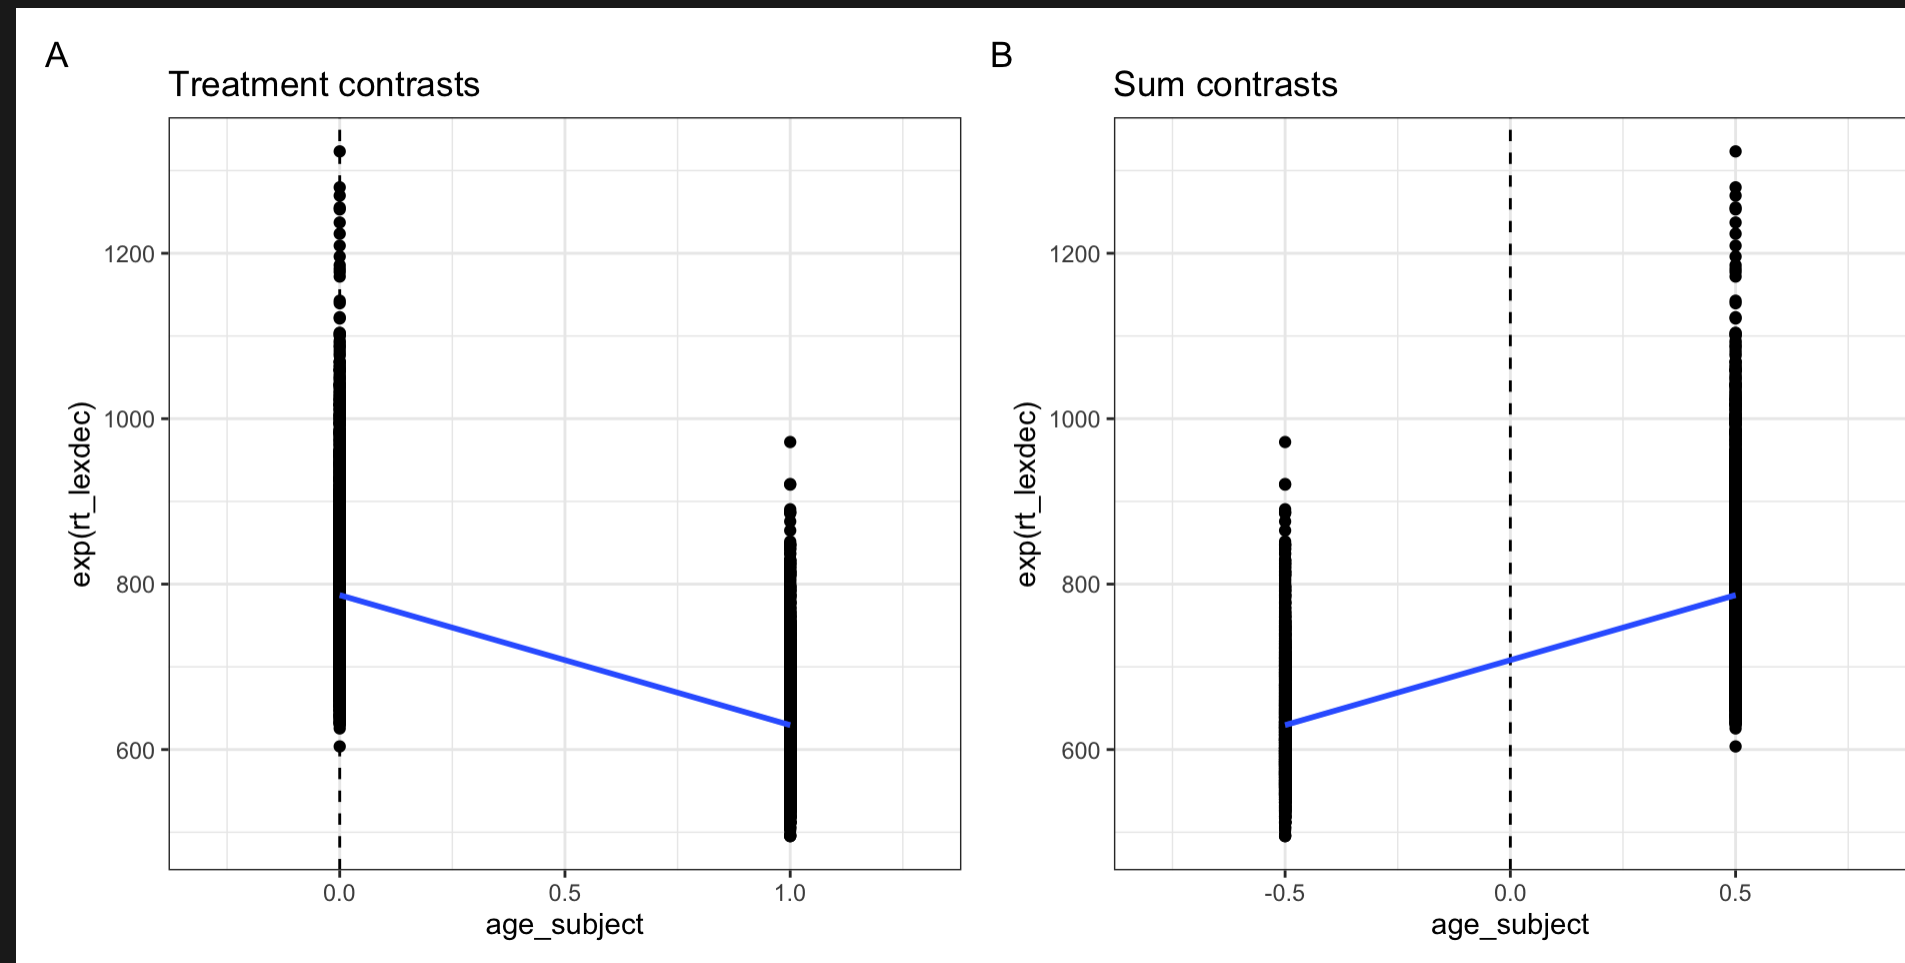
	term	estimate
	<chr>	<dbl>
1	(Intercept)	708.
2	age_subject1	157.

- there is a difference in the intercept
 - and a change in sign in our slope. Why is this?

Treatment/Dummy vs. Sum contrasts

► Code

Figure 1: The difference in slope corresponds to which level is coded as 0 (dummy coding) or -5/-1 (sum coding)



Intercept

- the intercept value is now the overall mean of all observed reaction times, because now the y value when x equals zero lies in the middle of the two groups
- the slope magnitude (i.e., size of the value) hasn't changed, because the difference between the two group means has not changed

```
1 mean(exp(df_freq_eng$rt_lexdec))  
[1] 708.1336
```

Exploring predicted values

- let's explore the predicted values of our model with a categorical variable

```
1 fitted(fit_age)[1:6]
      338      1790      3125      3957      3313      4145
629.5473 786.7200 629.5473 786.7200 629.5473 786.7200
```

- there are only 2 values, 630 and 787
 - these correspond to the means for each group that we saw above
 - they also seem to be in a pattern: mean(young), mean(old), mean(young), mean(old), etc.
 - how does this correspond to the age group of the participant for the first ten observations?

```
1 df_freq_eng$age_subject[1:6]
[1] young old  young old  young old
attr(,"contrasts")
      [,1]
old      0.5
young -0.5
Levels: old young
```


- first ten observations in our data are in young-old pairs. What are the first values in the raw data?

```
1 exp(df_freq_eng$rt_lexdec[1:6])
```

```
[1] 623.61 775.67 617.10 715.52 575.70 742.19
```

- what is the difference between these reaction times and the fitted values?

```
1 exp(df_freq_eng$rt_lexdec[1:6]) - fitted(fit_age)[1:6]
```

```
      338      1790      3125      3957      3313      4145  
-5.937299 -11.049991 -12.447299 -71.199991 -53.847299 -44.529991
```

```
1 residuals(fit_age)[1:6]
```

```
      338      1790      3125      3957      3313      4145  
-5.937299 -11.049991 -12.447299 -71.199991 -53.847299 -44.529991
```

- we see again that predicted values correspond to the x value for the corresponding row in the dataframe
 - but with our two-level factor, we only have two x values, young and old

```
1 df_freq_eng <-  
2   augment(fit_age, df_freq_eng)
```

```
1 df_freq_eng |>  
2   select(word, age_subject, rt_lexdec, .fitted, .resid) |>  
3   mutate(rt_lexdec = exp(rt_lexdec)) |>  
4   head()
```

A tibble: 6 × 5

	word	age_subject	rt_lexdec	.fitted	.resid
	<fct>	<fct>	<dbl>	<dbl>	<dbl>
1	ace	young	624.	630.	-5.94
2	ace	old	776.	787.	-11.0
3	act	young	617.	630.	-12.4
4	act	old	716.	787.	-71.2
5	add	young	576.	630.	-53.8
6	add	old	742.	787.	-44.5

Summary

- we saw that the equation for a straight line boils down to its intercept and slope
- we fit our first linear model with a categorical predictor

Important terms

term	description/other terms
------	-------------------------

Learning Objectives

Today we learned...

- about categorical predictors
- how to interpret different contrast coding

Task

Follow the instructions on the website ([Multiple regression > Task](#)) (or continue to the next slides).

Reading time data

We'll use a dataset from Biondo et al. (2022), an eye-tracking reading study exploring the processing of adverb-tense concord in Spanish past and future tenses. Participants read sentences that began with a temporal adverb (e.g., yesterday/tomorrow), and had a verb marked with the congruent or incongruent tense (past/future).

Load in the data.

```
1 df_tense <-  
2   read_csv(here("data", "Biondo.Soilemezidi.Mancini_dataset_ET.csv"),  
3             locale = locale(encoding = "Latin1") # for special characters in Spanish  
4             ) |>  
5   mutate(gramm = ifelse(gramm == "0", "ungramm", "gramm")) |>  
6   clean_names()
```


Treatment contrasts

We will look at the measure *total reading time* (**tt**) at the *verb* region (**roi == 4**). Subset the data to only include the verb region.

```
1 df_verb <-  
2   df_tense |>  
3   filter(roi == 4)
```

1. Run a simple linear model with (log-transformed) total reading time (**tt**) as an independent variable and grammaticality (**gramm**) as a dependent variable. Use treatment contrasts.
2. Inspect your coefficients again. What conclusions do you draw?
3. Run model diagnostics:
 - check model assumptions where relevant (normality, constant variance, collinearity)
 - check model fit (R^2)

Sum contrasts

1. Re-run your model with sum contrasts.
2. Inspect your coefficients again. Do your conclusions change?
3. Re-run your model diagnostics. How does it compare to your first model?

Multiple regression

1. Add verb tense (**verb_t**: past, future) as a predictor, including an interaction term. Use sum contrasts.
2. Inspect your coefficients again. Do your conclusions change?
3. Re-run your model diagnostics. How does it compare to the last models?

Literaturverzeichnis

Biondo, N., Soilemezidi, M., & Mancini, S. (2022). Yesterday is history, tomorrow is a mystery: An eye-tracking investigation of the processing of past and future time reference during sentence reading. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 48(7), 1001–1018. <https://doi.org/10.1037/xlm0001053>

