

Model selection

Parsimonious model selection

Daniela Palleschi

2024-02-09

Table of contents

Learning Objectives

Today we will...

- apply remedies for nonconvergence
- reduce our RES with a data-driven approach
- compare a parsimonious model to maximal and intercept-only models

Resources

- this lecture covers
 - Sections 10.3-5 in Sonderegger (2023)
 - Section 15.7.3 ‘Convergence Issues’ in Winter (2019)
 - Brauer & Curtin (2018)
 - Meteyard & Davies (2020)
- we will continue using the data from Biondo et al. (2022)

Set-up

```
# suppress scientific notation
options(scipen=999)

library(broman)
# function to format p-values
format_pval <- function(pval){
  dplyr::case_when(
    pval < .001 ~ "< .001",
    pval < .01 ~ "< .01",
    pval < .05 ~ "< .05",
    TRUE ~ broman::myround(pval, 3)
  )
}
```

Load packages

```
# load libraries
pacman::p_load(
  tidyverse,
  here,
  janitor,
  # new packages for mixed models:
  lme4,
  lmerTest,
  broom.mixed,
  lattice)

lmer <- lmerTest::lmer
```

Load data

- data from Biondo et al. (2022)

```
df_biondo <-
  read_csv(here("data", "Biondo.Soilemezidi.Mancini_dataset_ET.csv"),
    locale = locale(encoding = "Latin1") ## for special characters in Spanish
```

```

    ) |>
  clean_names() |>
  mutate(gramm = ifelse(gramm == "0", "ungramm", "gramm")) |>
  mutate_if(is.character, as_factor) |> # all character variables as factors
  droplevels() |>
  filter(adv_type == "Deic")

```

Set contrasts

```

contrasts(df_biondo$verb_t) <- c(-0.5, +0.5)
contrasts(df_biondo$gramm) <- c(-0.5, +0.5)

```

```

contrasts(df_biondo$verb_t)

```

```

      [,1]
Past    -0.5
Future   0.5

```

```

contrasts(df_biondo$gramm)

```

```

      [,1]
gramm    -0.5
ungramm   0.5

```

Start maximal

- model structure should be decided *a priori*
 - included fixed (predictors and covariates) and random effects

Maximal model

- starting point: most maximal model structure justified by your design
 - if this converges, great!
 - if it doesn't, what does this mean and what should we do?

```
fit_verb_fp_mm <- lmer(log(fp) ~ verb_t*gramm +
  (1 + verb_t*gramm|sj) +
  (1 + verb_t*gramm|item),
  data = df_biondo,
  subset = roi == 4)
```

- we get a warning of singular fit

Convergence issues

- “*Convergence is not a metric of model quality*” (Sonderegger, 2023, p. 365, Box 10.2)
 - convergence does not always indicate “overfitting” or “overparameterisation”
 - can also be due to optimizer choice
 - * since default optimizer was changed to `nloptwrap` from `bobyqa`, there seem to be more ‘false positive’ convergence warnings
- false-positive convergence: you get a convergence warning, but changing the optimizer and/or iteration count does not produce a warning
- false-negative convergence: you do not get a warning, but your variance-covariance matrix might indicate overfitting

Nonconvergence remedies

- unfortunately there is no one “right” way to deal with convergence issues
 - important is to transparently report and justify your method
- Table 17 in Brauer & Curtin (2018) (p. 404) suggests 20 remedies, whittled down to 10 suggestions in Sonderegger (2023)

Intrusive vs. Non-intrusive remedies

- non-intrusive remedies amount to checking/adjusting data and model specifications
- intrusive remedies involve reducing random effects structure
 - there are different schools of thought
 - * random-intercepts only: increased Type I error rate = overconfident estimates
 - * maximal-but-singular-fit model (Barr et al., 2013): reduces power = underconfident estimates

Table 10.1

Possible fixes for non-convergent (non-intrusive + intrusive) and singular models (intrusive only), ordered by which to try first (adapted from Brauer and Curtin 2018). Fixes 2(a) and 2(b) are tied.

1. Nonintrusive
 - a. Check your data and model
 - b. Standardize predictors (center, possibly scale)
 - c. Increase number of iterations
 - d. Change the optimizer
 - e. Give the optimizer better start values
2. Intrusive
 - a. Remove random effects involving control predictors (must not be in interactions with critical predictors)
 - b. Selectively remove random-effect correlations: for control predictors, then correlations that are probably close to 0
 - c. Remove random intercept (leaving slope terms in)
 - d. Remove random slopes for critical predictors

Figure 1: From Sonderegger (2023), p. 366

* data-driven approach (Bates et al., 2015): can lose the forest for the trees, e.g., removing random slopes for predictors of interest

- each strategy has its drawback
 - important is to choose your strategy *a priori* and transparently report and justify your strategy
 - even better: share/publish your data and code, which should be reproducible

?convergence

- type ?convergence in the Console and read the vignette
 - what suggestions does it make?
- compare this to ?isSingular

Non-intrusive methods

- check your data structure/variables
 - check model assumptions (e.g., normality, missing transformations of variables)
 - check your RES is justified by your experimental design/data structure

- centre your predictors (e.g., sum contrasts, or centring/standardizing) to reduce multicollinearity; reduces collinearity in the random effects (a possible source of nonconvergence)
- check observations per cell (e.g., is there a participant with very few observations, or few observations per one condition? Should be at least >5 per cell)
- alter model controls:
 - increase iterations
 - check optimizer

Check optimizer

- optimizer
 - `lme4::allFit(model)` (can take a while to run)

```
all_fit_verb_fp_mm <- allFit(fit_verb_fp_mm)
# bobyqa : boundary (singular) fit: see help('isSingular')
# [OK]
# Nelder_Mead : [OK]
# nlminbwrap : boundary (singular) fit: see help('isSingular')
# [OK]
# nmkbw : [OK]
# optimx.L-BFGS-B : boundary (singular) fit: see help('isSingular')
# [OK]
# nloptwrap.NLOPT_LN_NELDERMEAD : boundary (singular) fit: see help('isSingular')
# [OK]
# nloptwrap.NLOPT_LN_BOBYQA : boundary (singular) fit: see help('isSingular')
# [OK]
# There were 11 warnings (use warnings() to see them)
```

Optimizers

- default optimizer for `lmer()` is `nloptwrap`, formerly `bobyqa` (Bound Optimization by Quadratic Approximation)
 - usually changing to `bobyqa` helps
- see `?lmerControl` for more info
- if fits are very similar (or all optimizers except the default), the nonconvergent fit was a false positive

- it's safe to use the new optimizer

```
summary(all_fit_verb_fp_mm)$llik
```

bobyqa	Nelder_Mead
-2105.109	-2179.479
nlminbwrap	nmkbw
-2105.106	-2105.109
optimx.L-BFGS-B	nloptwrap.NLOPT_LN_NELDERMEAD
-2105.106	-2105.106
nloptwrap.NLOPT_LN_BOBYQA	
-2105.106	

```
summary(all_fit_verb_fp_mm)$fixef
```

	(Intercept)	verb_t1	gramm1	verb_t1:gramm1
bobyqa	5.956403	0.06170602	0.003369634	-0.01418865
Nelder_Mead	5.956350	0.06188102	0.003488675	-0.01397531
nlminbwrap	5.956403	0.06170726	0.003369637	-0.01419047
nmkbw	5.956404	0.06170653	0.003369153	-0.01419036
optimx.L-BFGS-B	5.956403	0.06170717	0.003369787	-0.01419044
nloptwrap.NLOPT_LN_NELDERMEAD	5.956403	0.06170725	0.003369649	-0.01419046
nloptwrap.NLOPT_LN_BOBYQA	5.956403	0.06170771	0.003369203	-0.01419184

Increase iterations

- and/or increase number of iterations
 - default is 10 000 (1e5 in scientific notation)
 - you can try 20 000, 100 000, etc.
 - this usually helps with larger data or models with complex RES

```
# check n of iterations
fit_verb_fp_mm@optinfo$feval
```

```
[1] 2318
```

`lmerControl()`

```
fit_verb_fp_mm <- lmer(log(fp) ~ verb_t*gramm +
  (1 + verb_t*gramm|sj) +
  (1 + verb_t*gramm|item),
  data = df_biondo,
  subset = roi == 4,
  control = lmerControl(optimizer = "bobyqa",
    optCtrl = list(maxfun = 2e5))
)
```

- or you can just ‘update’ the model to save some syntax

```
fit_verb_fp_mm <- update(fit_verb_fp_mm,
  control = lmerControl(optimizer = "bobyqa",
    optCtrl = list(maxfun = 2e5)))
```

boundary (singular) fit: see `help('isSingular')`

Warning: Model failed to converge with 1 negative eigenvalue: -5.3e-01

Removing parameters

- still won't converge?
 - it's time to consider intrusive remedies: removing random effects parameters

Intrusive methods

- nonconvergence in maximal models is often due to overfitting
 - i.e., the model is overly complex given your data
 - this is typically due to an overly complex random effects structure
- if the non-intrusive methods don't lead to convergence, the problem is likely overfitting

Parsimonious vs. maximal

- there are different camps on how to deal with this issue
- I personally follow the suggestions in Bates et al. (2015) (for now)
 1. run random effects Principal Components Analysis (`summary(rePCA(model))`, `lme4` package)
 - informs by how many parameters our model is overfit
 2. check variance-covariance matrix (`VarCorr(model)`)
 3. remove parameters with very high or low Correlation terms and/or much lower variance compared to other terms
 4. fit simplified model
 5. wash, rinse, repeat
- we'll practice this method today, but keep in mind that it's up to you to decide and justify which method you use

Random effects Principal Components Analysis

- gives us a ranking of all parameters ('components') in our RES per unit

```
summary(rePCA(fit_verb_fp_mm))
```

\$item

Importance of components:

	[,1]	[,2]	[,3]	[,4]
Standard deviation	0.3638	0.2493	0.08366	0.000000000000000004965
Proportion of Variance	0.6567	0.3085	0.03474	0.000000000000000000000
Cumulative Proportion	0.6567	0.9653	1.00000	1.000000000000000000000

\$sj

Importance of components:

	[,1]	[,2]	[,3]	[,4]
Standard deviation	0.6490	0.01470	0.000007463	0.0000001104
Proportion of Variance	0.9995	0.00051	0.000000000	0.0000000000
Cumulative Proportion	0.9995	1.00000	1.000000000	1.0000000000

- important is the Cumulative Proportion

- how much of the cumulative variance explained by all the by-unit parameters does this one parameter contribute?
- we see for item, the first component accounts for 66% of the variance explained, and the next contributes an additional 31%, and the next 3%
- so two components account for roughly 97% of variance explained by our RES
- in other words, we can remove one component for sure, and possibly another
- we could potentially remove 3 components from participant

Variance-covariance matrix

- so we can remove 2 parameters from item and participant
 - so either the varying intercept, or slope for tense, grammaticality, or their interaction
- we can check this with `VarCorr(fit_verb_fp_mm)`

```
VarCorr(fit_verb_fp_mm)
```

Groups	Name	Std.Dev.	Corr			
item	(Intercept)	0.139189				
	verb_t1	0.055890	0.488			
	gramm1	0.022569	-0.109	-0.921		
	verb_t1:gramm1	0.095313	-0.283	0.456	-0.646	
sj	(Intercept)	0.257535				
	verb_t1	0.018297	0.974			
	gramm1	0.012055	0.960	0.872		
	verb_t1:gramm1	0.017731	0.990	0.933	0.990	
Residual		0.399095				

- for item I would remove **gramm** because it has the lowest variance, and has a pretty high correlation with **verb_t** (which is unlikely to be true)
- I would also remove **gramm** for participant for the same reason, as well as its high correlation with the intercept and **verb_t**

Alternate model 1

- for now let's just remove the interaction term
 - for reproducibility reasons, do not delete the code for a model that did not converge
 - rather, write a comment on what decision was made (and why) for the new model

```

fit_verb_fp_m1 <- lmer(log(fp) ~ verb_t*gramm +
  (1 + verb_t+gramm|sj) +
  (1 + verb_t+gramm|item),
  data = df_biondo,
  subset = roi == 4,
  control = lmerControl(optimizer = "bobyqa",
    optCtrl = list(maxfun = 2e5))
)

```

boundary (singular) fit: see help('isSingular')

rePCA()

```
summary(rePCA(fit_verb_fp_m1))
```

\$item

Importance of components:

	[,1]	[,2]	[,3]
Standard deviation	0.3559	0.1291	0
Proportion of Variance	0.8837	0.1163	0
Cumulative Proportion	0.8837	1.0000	1

\$sj

Importance of components:

	[,1]	[,2]	[,3]
Standard deviation	0.6465	0.0000004537	0
Proportion of Variance	1.0000	0.0000000000	0
Cumulative Proportion	1.0000	1.0000000000	1

VarCorr()

```
VarCorr(fit_verb_fp_m1)
```

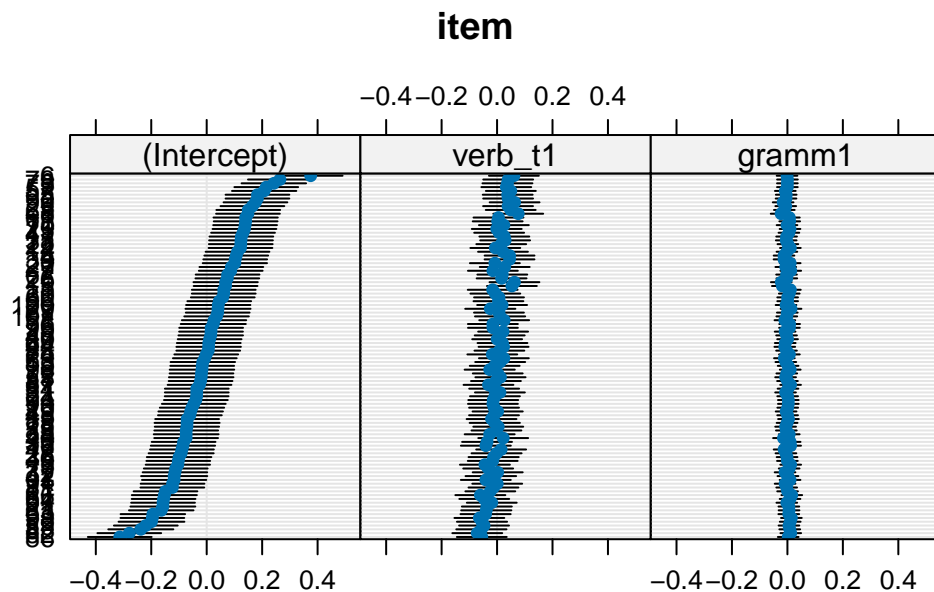
Groups	Name	Std.Dev.	Corr
item	(Intercept)	0.139274	
	verb_t1	0.055550	0.489
	gramm1	0.020747	-0.117 -0.924
sj	(Intercept)	0.257657	
	verb_t1	0.017584	1.000

	gramm1	0.011554	1.000	1.000
Residual		0.399869		

- when we see Corr ± 1 , this tells us there was an error computing correlations between parameters
 - it is an invitation to explore
- this is not plausible, and indicates overfitting in our model
 - we can remove all slopes from sj


by-item random effects

```
lattice::dotplot(ranef(fit_verb_fp_m1))$item
```



by-participant random effects (with +1 correlations)

```
lattice::dotplot(ranef(fit_verb_fp_m1))$sj
```



Alternate model 2

```
fit_verb_fp_m2 <- lmer(log(fp) ~ verb_t*gramm +
  (1 |sj) +
  (1 + verb_t+gramm|item),
  data = df_biondo,
  subset = roi == 4,
  control = lmerControl(optimizer = "bobyqa",
    optCtrl = list(maxfun = 2e5))
)
```

boundary (singular) fit: see help('isSingular')

rePCA()

```
summary(rePCA(fit_verb_fp_m2))
```

\$item

Importance of components:

	[,1]	[,2]	[,3]
Standard deviation	0.3559	0.1297	0
Proportion of Variance	0.8827	0.1173	0

```
Cumulative Proportion 0.8827 1.0000 1
```

```
$sj
```

```
Importance of components:
```

```
[,1]
```

```
Standard deviation 0.6441
```

```
Proportion of Variance 1.0000
```

```
Cumulative Proportion 1.0000
```

```
VarCorr()
```


```
VarCorr(fit_verb_fp_m2)
```

Groups	Name	Std.Dev.	Corr
item	(Intercept)	0.139364	
	verb_t1	0.055805	0.485
	gramm1	0.020546	-0.097 -0.917
sj	(Intercept)	0.257648	
Residual		0.399995	

- by-item slopes for `gramm` and `verb_t` are highly correlated
- `gramm` has least variance, so let's remove it

by-item random effects

```
lattice::dotplot(ranef(fit_verb_fp_m2))$item
```



Alternate model 3

```
fit_verb_fp_m3 <- lmer(log(fp) ~ verb_t*gramm +
  (1 |sj) +
  (1 + verb_t|item),
  data = df_biondo,
  subset = roi == 4,
  control = lmerControl(optimizer = "bobyqa",
    optCtrl = list(maxfun = 2e5))
)
```

- converged!

rePCA()

```
summary(rePCA(fit_verb_fp_m3))
```

\$item

Importance of components:

	[,1]	[,2]
Standard deviation	0.3553	0.10311
Proportion of Variance	0.9223	0.07768
Cumulative Proportion	0.9223	1.00000

```
$sj
Importance of components:
              [,1]
Standard deviation    0.6438
Proportion of Variance 1.0000
Cumulative Proportion 1.0000
```

```
VarCorr()
```

```
VarCorr(fit_verb_fp_m3)
```

Groups	Name	Std.Dev.	Corr
item	(Intercept)	0.139365	
	verb_t1	0.050134	0.542
sj	(Intercept)	0.257714	
Residual		0.400315	

Alternate model 4

- but we might've also decided to remove `verb_t`, so let's run that model

```
fit_verb_fp_m4 <- lmer(log(fp) ~ verb_t*gramm +
                      (1 |sj) +
                      (1 + gramm|item),
                      data = df_biondo,
                      subset = roi == 4,
                      control = lmerControl(optimizer = "bobyqa",
                                             optCtrl = list(maxfun = 2e5))
                      )
```

boundary (singular) fit: see `help('isSingular')`

- does not converge, so we're justified in keeping by-item `verb_t` slopes


Final model

- the final model name should be some sort of convention to make your life easier
 - so remove model index


```
fit_verb_fp <- fit_verb_fp_m3
```

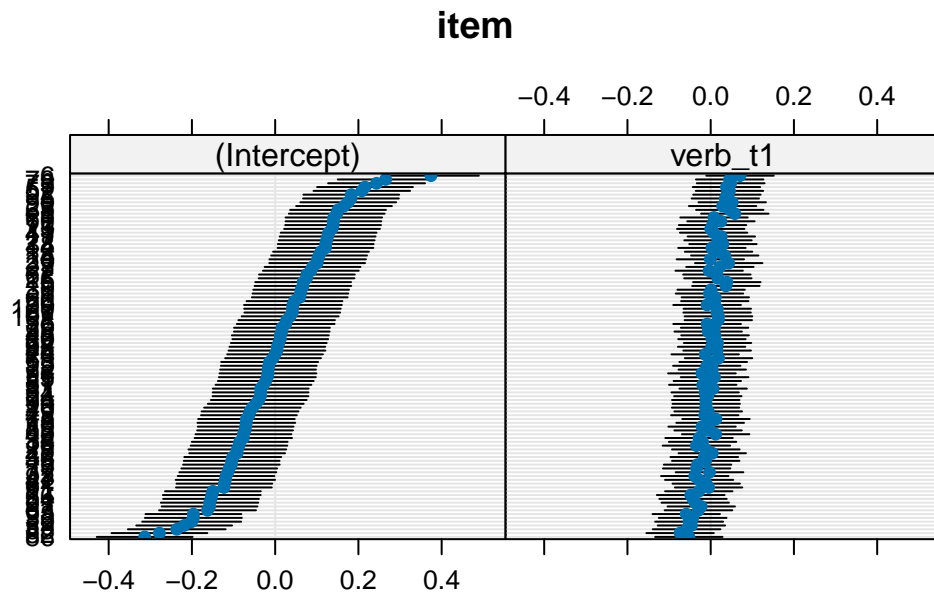
by-item random effects

```
lattice::dotplot(ranef(fit_verb_fp))$sj
```



by-participant random effects

```
lattice::dotplot(ranef(fit_verb_fp))$item
```



```
summary()
```

```
summary(fit_verb_fp)
```

```
Linear mixed model fit by REML. t-tests use Satterthwaite's method [
lmerModLmerTest]
Formula: log(fp) ~ verb_t * gramm + (1 | sj) + (1 + verb_t | item)
  Data: df_biondo
Control: lmerControl(optimizer = "bobyqa", optCtrl = list(maxfun = 200000))
Subset: roi == 4
```

```
REML criterion at convergence: 4216.2
```

```
Scaled residuals:
```

```
      Min       1Q   Median       3Q      Max
-4.1758 -0.6096 -0.0227  0.6060  4.0568
```

```
Random effects:
```

```
Groups   Name             Variance Std.Dev. Corr
item     (Intercept) 0.019423 0.13936
         verb_t1      0.002513 0.05013  0.54
sj       (Intercept) 0.066417 0.25771
Residual                0.160252 0.40032
Number of obs: 3795, groups: item, 96; sj, 60
```

Fixed effects:

	Estimate	Std. Error	df	t value	Pr(> t)
(Intercept)	5.956384	0.036763	79.243183	162.021	< 0.00000000000000002
verb_t1	0.061733	0.013971	93.410602	4.419	0.0000267
gramm1	0.003298	0.012999	3544.451823	0.254	0.80
verb_t1:gramm1	-0.014380	0.025998	3544.762347	-0.553	0.58

(Intercept) ***

verb_t1 ***

gramm1

verb_t1:gramm1

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

	(Intr)	vr_b_t1	gramm1
verb_t1		0.077	
gramm1	0.000	-0.002	
vr_b_t1:grm1	0.000	0.002	0.000

- IMPORTANTLY, only look at the fixed effects after you've got your final model!!!!
 - i.e., run model -> convergence error -> `rePCA()` + `VarCorr()` -> run model -> ...
 - > converges -> only NOW run `summary(model)`

Comparing to 'bad' models

- let's compare our final model to our 'bad' models
 - random intercepts-only model (overconfident)
 - maximal model (underconfident)

Random-intercepts only

```
fit_verb_fp_intercepts <- lmer(log(fp) ~ verb_t*gramm +
  (1 |sj) +
  (1 |item),
  data = df_biondo,
  subset = roi == 4
)
```