

Model selection

Parsimonious model selection

Daniela Palleschi

Humboldt-Universität zu Berlin

2024-02-09

Learning Objectives

Today we will...

- apply remedies for nonconvergence
- reduce our RES with a data-driven approach
- compare a parsimonious model to maximal and intercept-only models

Resources

- this lecture covers
 - Sections 10.3-5 in Sonderegger ([2023](#))
 - Section 15.7.3 ‘Convergence Issues’ in Winter ([2019](#))
 - Brauer & Curtin ([2018](#))
 - Meteyard & Davies ([2020](#))
- we will continue using the data from Biondo et al. ([2022](#))

Set-up

```
1 # suppress scientific notation  
2 options(scipen=999)
```

► Code for a function to format p-values

Load packages

```
1 # load libraries
2 pacman::p_load(
3     tidyverse,
4     here,
5     janitor,
6     # new packages for mixed models:
7     lme4,
8     lmerTest,
9     broom.mixed,
10    lattice)
```

```
1 lmer <- lmerTest::lmer
```

Load data

- data from Biondo et al. ([2022](#))

```
1 df_biondo <-  
2   read_csv(here("data", "Biondo.Soilemezidi.Mancini_dataset_ET.csv"),  
3            locale = locale(encoding = "Latin1") ## for special characters in Spanish  
4            ) |>  
5   clean_names() |>  
6   mutate(gramm = ifelse(gramm == "0", "ungramm", "gramm")) |>  
7   mutate_if(is.character, as_factor) |> # all character variables as factors  
8   droplevels() |>  
9   filter(adv_type == "Deic")
```

Set contrasts

```
1 contrasts(df_biondo$verb_t) <- c(-0.5,+0.5)
2 contrasts(df_biondo$gramm) <- c(-0.5,+0.5)
```

```
1 contrasts(df_biondo$verb_t)
```

```
      [,1]
Past    -0.5
Future   0.5
```

```
1 contrasts(df_biondo$gramm)
```

```
      [,1]
gramm    -0.5
ungramm   0.5
```

Start maximal

- model structure should be decided *a priori*
 - included fixed (predictors and covariates) and random effects

Maximal model

- starting point: most maximal model structure justified by your design
 - if this converges, great!
 - if it doesn't, what does this mean and what should we do?

```
1 fit_verb_fp_mm <- lmer(log(fp) ~ verb_t*gramm +  
2                       (1 + verb_t*gramm|sj) +  
3                       (1 + verb_t*gramm|item),  
4                       data = df_biondo,  
5                       subset = roi == 4)
```

- we get a warning of singular fit

Convergence issues

- “*Convergence is not a metric of model quality*” ([Sonderegger, 2023, p. 365](#), Box 10.2)
 - convergence does not always indicate “overfitting” or “overparameterisation”
 - can also be due to optimizer choice
 - since default optimizer was changed to **nloptwrap** from **bobyqa**, there seem to be more ‘false positive’ convergence warnings
- false-positive convergence: you get a convergence warning, but changing the optimizer and/or iteration count does not produce a warning
- false-negative convergence: you do not get a warning, but your variance-covariance matrix might indicate overfitting

Nonconvergence remedies

- unfortunately there is no one “right” way to deal with convergence issues
 - important is to transparently report and justify your method
- Table 17 in Brauer & Curtin (2018) (p. 404) suggests 20 remedies, whittled down to 10 suggestions in Sonderegger (2023)

Table 10.1

Possible fixes for non-convergent (non-intrusive + intrusive) and singular models (intrusive only), ordered by which to try first (adapted from Brauer and Curtin 2018). Fixes 2(a) and 2(b) are tied.

1. Nonintrusive

- a. Check your data and model
- b. Standardize predictors (center, possibly scale)
- c. Increase number of iterations
- d. Change the optimizer
- e. Give the optimizer better start values

2. Intrusive

- a. Remove random effects involving control predictors (must not be in interactions with critical predictors)
- b. Selectively remove random-effect correlations: for control predictors, then correlations that are probably close to 0
- c. Remove random intercept (leaving slope terms in)
- d. Remove random slopes for critical predictors

Intrusive vs. Non-intrusive remedies

?convergence

- type ?convergence in the Console and read the vignette
 - what suggestions does it make?
- compare this to ?isSingular

Non-intrusive methods

- check your data structure/variables
 - check model assumptions (e.g., normality, missing transformations of variables)
 - check your RES is justified by your experimental design/data structure
 - centre your predictors (e.g., sum contrasts, or centring/standardizing) to reduce multicollinearity; reduces collinearity in the random effects (a possible source of nonconvergence)
 - check observations per cell (e.g., is there a participant very few observations, or few observations per one condition? Should be at least >5 per cell)
- alter model controls:
 - increase iterations
 - check optimizer

Check optimizer

- optimizer
 - `lme4::allFit(model)` (can take a while to run)

```
1 all_fit_verb_fp_mm <- allFit(fit_verb_fp_mm)
2 # bobyqa : boundary (singular) fit: see help('isSingular')
3 # [OK]
4 # Nelder_Mead : [OK]
5 # nlminbwrap : boundary (singular) fit: see help('isSingular')
6 # [OK]
7 # nmkbw : [OK]
8 # optimx.L-BFGS-B : boundary (singular) fit: see help('isSingular')
9 # [OK]
10 # nloptwrap.NLOPT_LN_NELDERMEAD : boundary (singular) fit: see help('isSingular')
11 # [OK]
12 # nloptwrap.NLOPT_LN_BOBYQA : boundary (singular) fit: see help('isSingular')
13 # [OK]
14 # There were 11 warnings (use warnings() to see them)
```

Optimizers

- default optimizer for `lmer()` is `nloptwrap`, formerly `bobyqa` (Bound Optimization by Quaradric Approximation)
 - usually changing to `bobyqa` helps
- see `?lmerControl` for more info
- if fits are very similar (or all optimizers except the default), the nonconvergent fit was a false positive
 - it's safe to use the new optimizer

```
1 summary(all_fit_verb_fp_mm)$llik
```

bobyqa	Nelder_Mead
-2105.109	-2179.479
nlminbwrap	nmkbw
-2105.106	-2105.109
optimx.L-BFGS-B	nloptwrap.NLOPT_LN_NELDERMEAD
-2105.106	-2105.106
nloptwrap.NLOPT_LN_BOBYQA	
-2105.106	

```
1 summary(all_fit_verb_fp_mm)$fixef
```

	(Intercept)	verb_t1	gramm1	verb_t1:gramm1
bobyqa	5.956403	0.06170602	0.003369634	-0.01418865
Nelder_Mead	5.956350	0.06188102	0.003488675	-0.01397531
nlminbwrap	5.956403	0.06170726	0.003369637	-0.01419047

nmkbw	5.956404	0.06170653	0.003369153	-0.01419036
optimx.L-BFGS-B	5.956403	0.06170717	0.003369787	-0.01419044
nloptwrap.NLOPT_LN_NELDERMEAD	5.956403	0.06170725	0.003369649	-0.01419046
nloptwrap.NLOPT_LN_BOBYQA	5.956403	0.06170771	0.003369203	-0.01419184

Increase iterations

- and/or increase number of iterations
 - default is 10 000 (**1e5** in scientific notation)
 - you can try 20 000, 100 000, etc.
 - this usually helps with larger data or models with complex RES

```
1 # check n of iterations
2 fit_verb_fp_mm@optinfo$feval
```

```
[1] 2318
```

lmerControl()

```
1 fit_verb_fp_mm <- lmer(log(fp) ~ verb_t*gramm +
2                       (1 + verb_t*gramm|sj) +
3                       (1 + verb_t*gramm|item),
4                       data = df_biondo,
5                       subset = roi == 4,
6                       control = lmerControl(optimizer = "bobyqa",
7                                             optCtrl = list(maxfun = 2e5))
8 )
```

- or you can just ‘update’ the model to save some syntax

```
1 fit_verb_fp_mm <- update(fit_verb_fp_mm,
2                          control = lmerControl(optimizer = "bobyqa",
3                                                optCtrl = list(maxfun = 2e5)))
```

boundary (singular) fit: see help('isSingular')

Warning: Model failed to converge with 1 negative eigenvalue: -5.3e-01

Removing parameters

- still won't converge?
 - it's time to consider intrusive remedies: removing random effects parameters

Intrusive methods

- nonconvergence in maximal models is often due to overfitting
 - i.e., the model is overly complex given your data
 - this is typically due to an overly complex random effects structure
- if the non-intrusive methods don't lead to convergence, the problem is likely overfitting

Parsimonious vs. maximal

- there are different camps on how to deal with this issue
- I personally follow the suggestions in Bates et al. (2015) (for now)
 1. run random effects Principal Components Analysis (`summary(rePCA(model))`, `lme4` package)
 - informs by how many parameters our model is overfit
 2. check variance-covariance matrix (`VarCorr(model)`)
 3. remove parameters with very high or low Correlation terms and/or much lower variance compared to other terms
 4. fit simplified model
 5. wash, rinse, repeat
- we'll practice this method today, but keep in mind that it's up to you to decide and justify which method you use

Random effects Principal Components Analysis

- gives us a ranking of all parameters (‘components’) in our RES per unit

```
1 summary(rePCA(fit_verb_fp_mm))
```

\$item

Importance of components:

	[,1]	[,2]	[,3]	[,4]
Standard deviation	0.3638	0.2493	0.08366	0.000000000000000000004965
Proportion of Variance	0.6567	0.3085	0.03474	0.000000000000000000000000
Cumulative Proportion	0.6567	0.9653	1.00000	1.000000000000000000000000

\$sj

Importance of components:

	[,1]	[,2]	[,3]	[,4]
Standard deviation	0.6490	0.01470	0.000007463	0.0000001104
Proportion of Variance	0.9995	0.00051	0.000000000	0.0000000000
Cumulative Proportion	0.9995	1.00000	1.000000000	1.0000000000

- important is the Cumulative Proportion
 - how much of the cumulative variance explained by all the by-unit parameters does this one parameter contribute?
 - we see for item, the first component accounts for 66% of the variance explained, and the next contributes an additional 31%, and the next 3%
 - so two components account for roughly 97% of variance explained by our RES
 - in other words, we can remove one component for sure, and possibly another
 - we could potentially remove 3 components from participant

Variance-covariance matrix

- so we can remove 2 parameters from item and participant
 - so either the varying intercept, or slope for tense, grammaticality, or their interaction
- we can check this with `VarCorr(fit_verb_fp_mm)`

```
1 VarCorr(fit_verb_fp_mm)
```

Groups	Name	Std.Dev.	Corr
item	(Intercept)	0.139189	
	verb_t1	0.055890	0.488
	gramm1	0.022569	-0.109 -0.921
	verb_t1:gramm1	0.095313	-0.283 0.456 -0.646
sj	(Intercept)	0.257535	
	verb_t1	0.018297	0.974
	gramm1	0.012055	0.960 0.872
	verb_t1:gramm1	0.017731	0.990 0.933 0.990
Residual		0.399095	

- for item I would remove `gramm` because it has the lowest variance, and has a pretty high correlation with `verb_t` (which is unlikely to be true)
- I would also remove `gramm` for participant for the same reason, as well as its high correlation with the intercept and `verb_t`

Alternate model 1

- for now let's just remove the interaction term
 - for reproducibility reasons, do not delete the code for a model that did not converge
 - rather, write a comment on what decision was made (and why) for the new model

```
1 fit_verb_fp_m1 <- lmer(log(fp) ~ verb_t*gramm +  
2                       (1 + verb_t+gramm|sj) +  
3                       (1 + verb_t+gramm|item),  
4                       data = df_biondo,  
5                       subset = roi == 4,  
6                       control = lmerControl(optimizer = "bobyqa",  
7                                             optCtrl = list(maxfun = 2e5))  
8 )
```

boundary (singular) fit: see help('isSingular')

rePCA()

```
1 summary(rePCA(fit_verb_fp_m1))
```

\$item

Importance of components:

	[,1]	[,2]	[,3]
Standard deviation	0.3559	0.1291	0
Proportion of Variance	0.8837	0.1163	0
Cumulative Proportion	0.8837	1.0000	1

\$sj

Importance of components:

	[,1]	[,2]	[,3]
Standard deviation	0.6465	0.0000004537	0
Proportion of Variance	1.0000	0.0000000000	0
Cumulative Proportion	1.0000	1.0000000000	1

VarCorr()

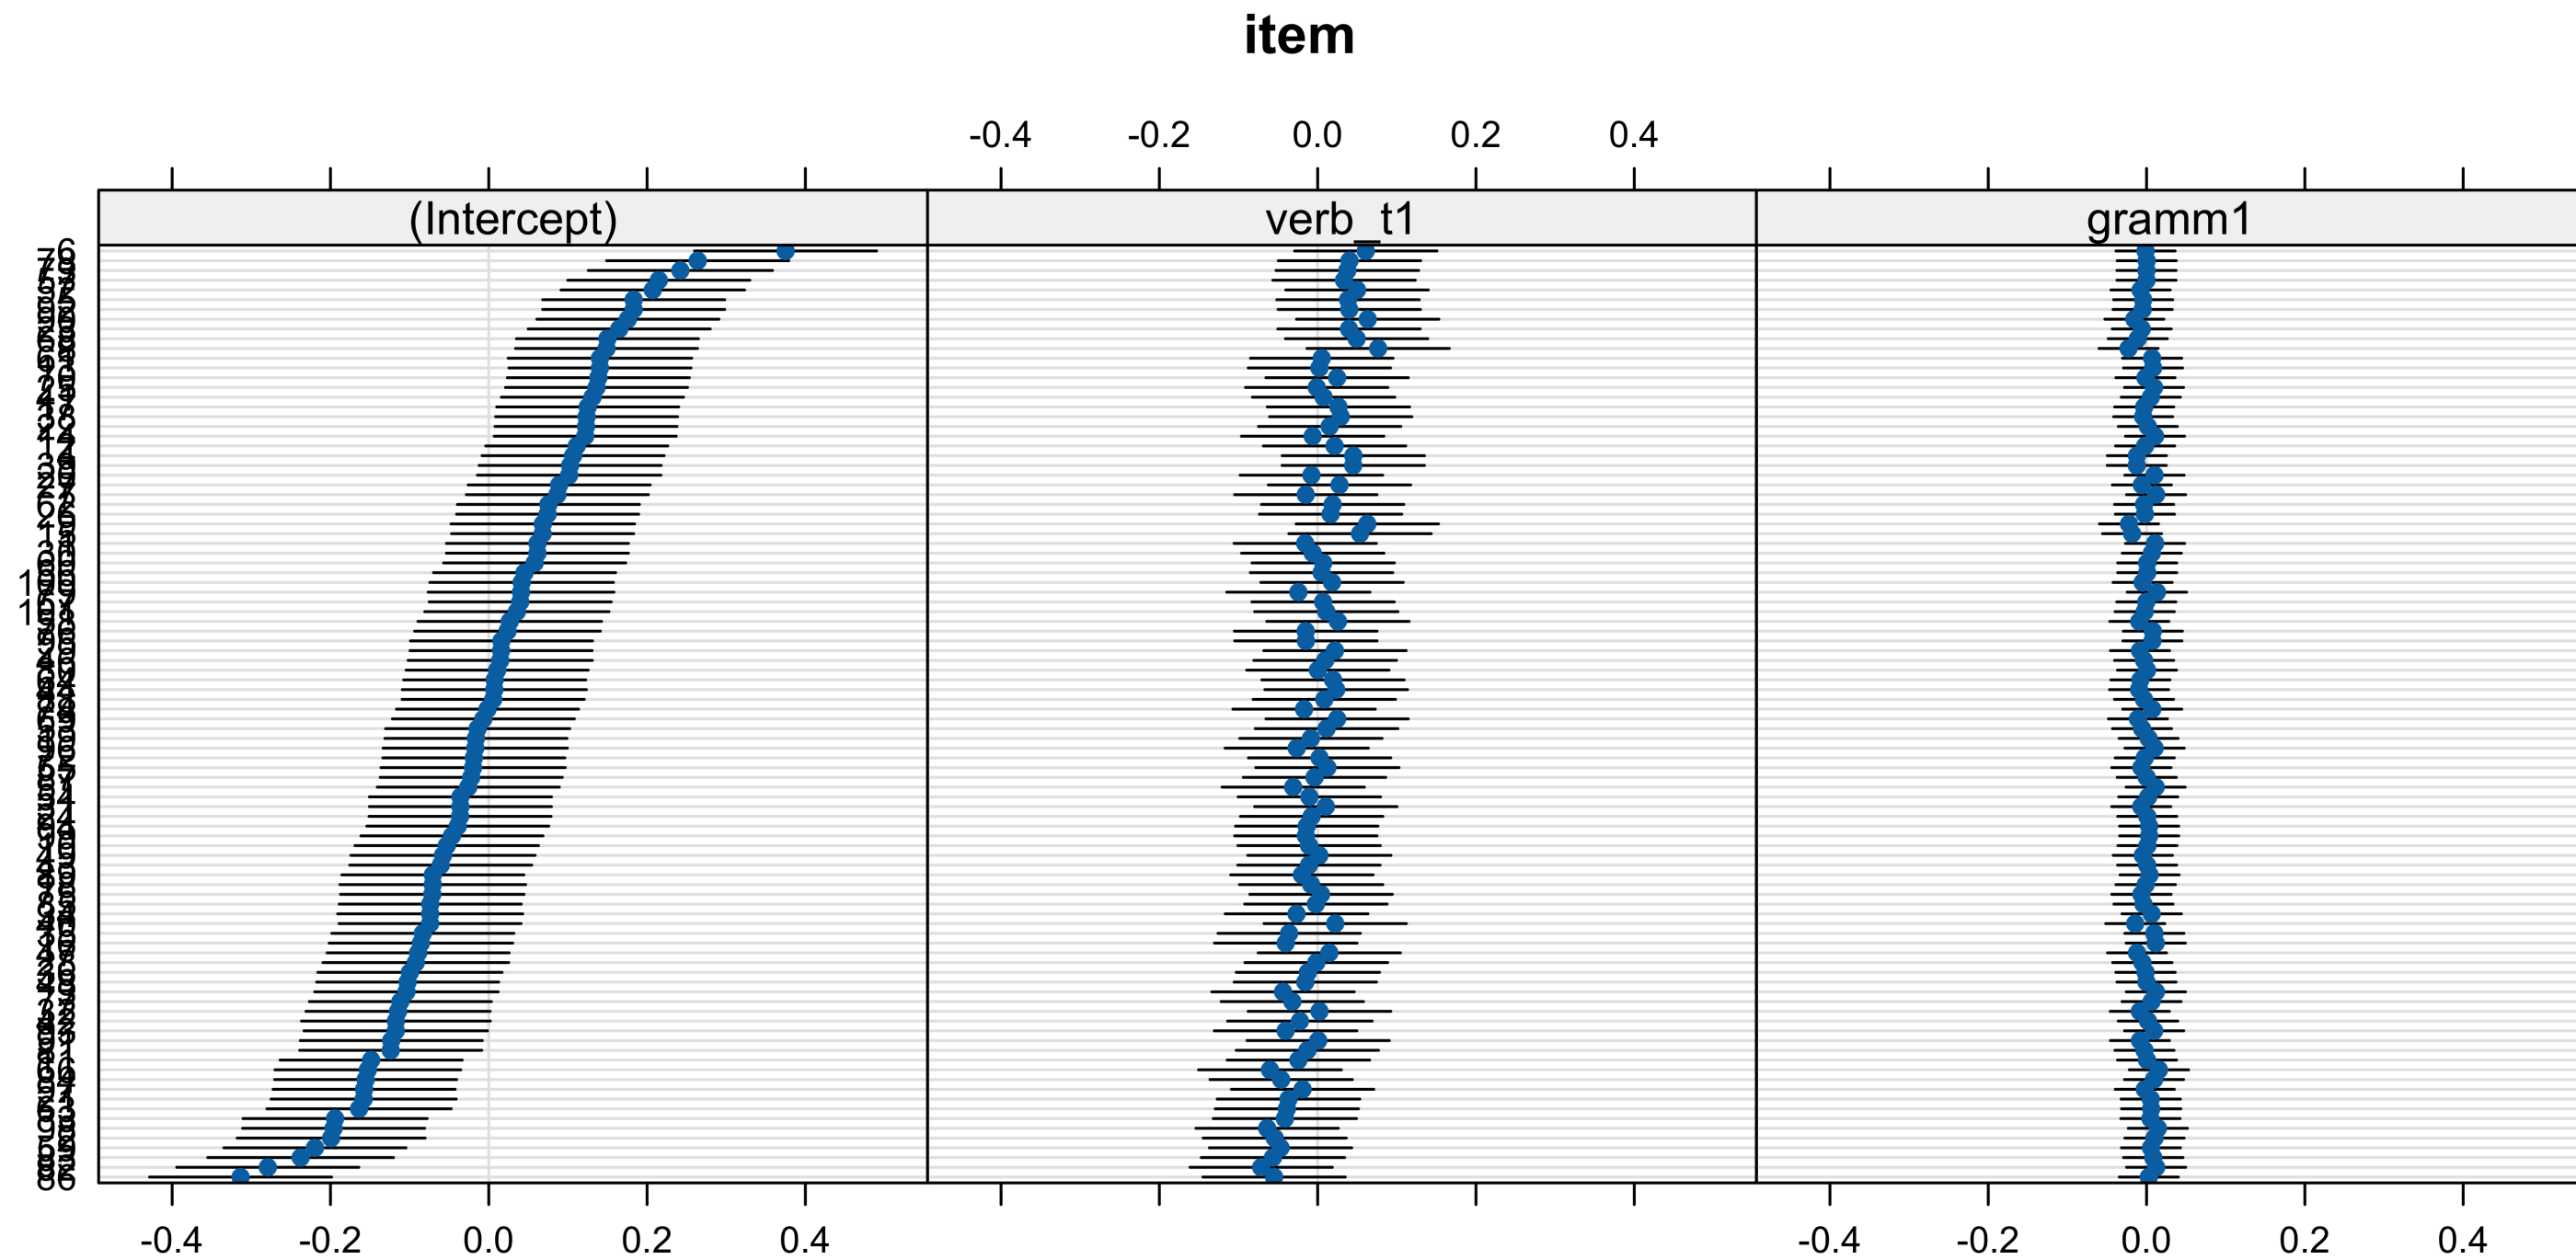
```
1 VarCorr(fit_verb_fp_m1)
```

Groups	Name	Std.Dev.	Corr
item	(Intercept)	0.139274	
	verb_t1	0.055550	0.489
	gramm1	0.020747	-0.117 -0.924
sj	(Intercept)	0.257657	
	verb_t1	0.017584	1.000
	gramm1	0.011554	1.000 1.000
Residual		0.399869	

- when we see Corr +/-1, this tells us there was an error computing correlations between parameters
 - it is an invitation to explore
- this is not plausible, and indicates overfitting in our model
 - we can remove all slopes from sj


by-item random effects

```
1 lattice::dotplot(ranef(fit_verb_fp_m1))$item
```



by-participant random effects (with +1 correlations)

```
1 lattice::dotplot(ranef(fit_verb_fp_m1))$sj
```



Alternate model 2

```
1 fit_verb_fp_m2 <- lmer(log(fp) ~ verb_t*gramm +  
2                       (1 |sj) +  
3                       (1 + verb_t+gramm|item),  
4                       data = df_biondo,  
5                       subset = roi == 4,  
6                       control = lmerControl(optimizer = "bobyqa",  
7                                             optCtrl = list(maxfun = 2e5))  
8 )
```

boundary (singular) fit: see help('isSingular')

rePCA()

```
1 summary(rePCA(fit_verb_fp_m2))
```

\$item

Importance of components:

	[,1]	[,2]	[,3]
Standard deviation	0.3559	0.1297	0
Proportion of Variance	0.8827	0.1173	0
Cumulative Proportion	0.8827	1.0000	1

\$sj

Importance of components:

	[,1]
Standard deviation	0.6441
Proportion of Variance	1.0000
Cumulative Proportion	1.0000

VarCorr()


```
1 VarCorr(fit_verb_fp_m2)
```

Groups	Name	Std.Dev.	Corr
item	(Intercept)	0.139364	
	verb_t1	0.055805	0.485
	gramm1	0.020546	-0.097 -0.917
sj	(Intercept)	0.257648	
Residual		0.399995	

- by-item slopes for **g**ramm and **v**erb_**t** are highly correlated
- **g**ramm has least variance, so let's remove it

by-item random effects

```
1 lattice::dotplot(ranef(fit_verb_fp_m2))$item
```



Alternate model 3

```
1 fit_verb_fp_m3 <- lmer(log(fp) ~ verb_t*gramm +  
2                       (1 |sj) +  
3                       (1 + verb_t|item),  
4                       data = df_biondo,  
5                       subset = roi == 4,  
6                       control = lmerControl(optimizer = "bobyqa",  
7                                             optCtrl = list(maxfun = 2e5))  
8 )
```

- converged!

rePCA()

```
1 summary(rePCA(fit_verb_fp_m3))
```

\$item

Importance of components:

	[,1]	[,2]
Standard deviation	0.3553	0.10311
Proportion of Variance	0.9223	0.07768
Cumulative Proportion	0.9223	1.00000

\$sj

Importance of components:

	[,1]
Standard deviation	0.6438
Proportion of Variance	1.0000
Cumulative Proportion	1.0000

VarCorr()

```
1 VarCorr(fit_verb_fp_m3)
```

Groups	Name	Std.Dev.	Corr
item	(Intercept)	0.139365	
	verb_t1	0.050134	0.542
sj	(Intercept)	0.257714	
Residual		0.400315	

Alternate model 4

- but we might've also decided to remove `verb_t`, so let's run that model

```
1 fit_verb_fp_m4 <- lmer(log(fp) ~ verb_t*gramm +  
2                       (1 |sj) +  
3                       (1 + gramm|item),  
4                       data = df_biondo,  
5                       subset = roi == 4,  
6                       control = lmerControl(optimizer = "bobyqa",  
7                                             optCtrl = list(maxfun = 2e5))  
8 )
```

boundary (singular) fit: see help('isSingular')

- does not converge, so we're justified in keeping by-item `verb_t` slopes

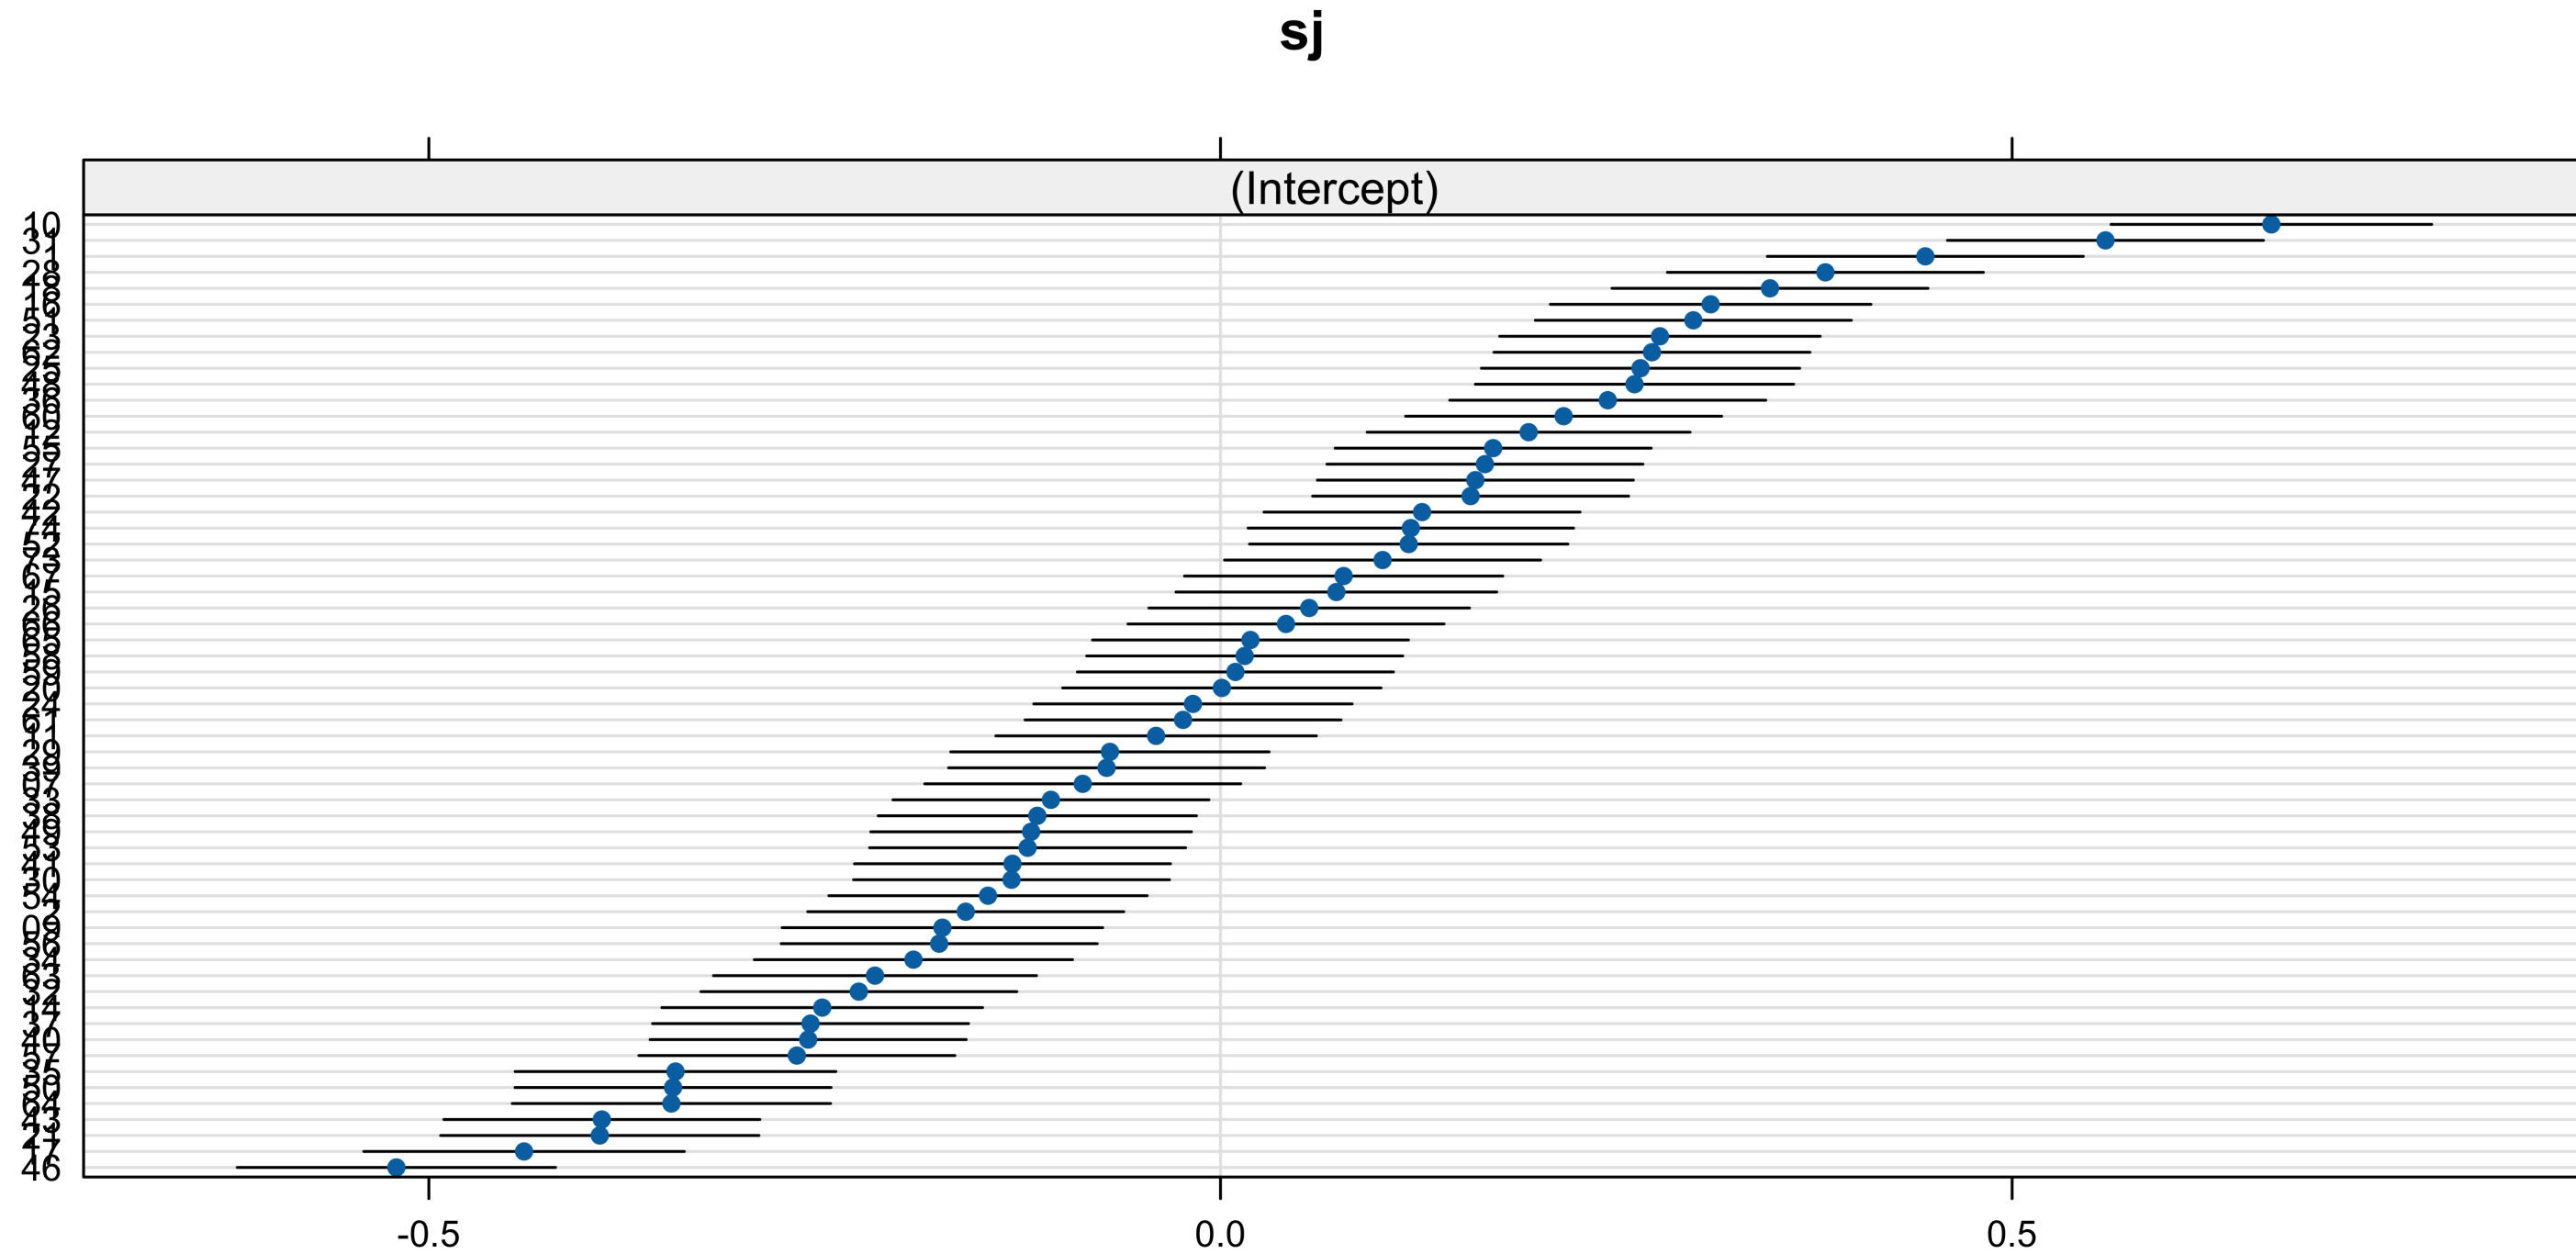
Final model

- the final model name should be some sort of convention to make your life easier
 - so remove model index

```
1 fit_verb_fp <- fit_verb_fp_m3
```

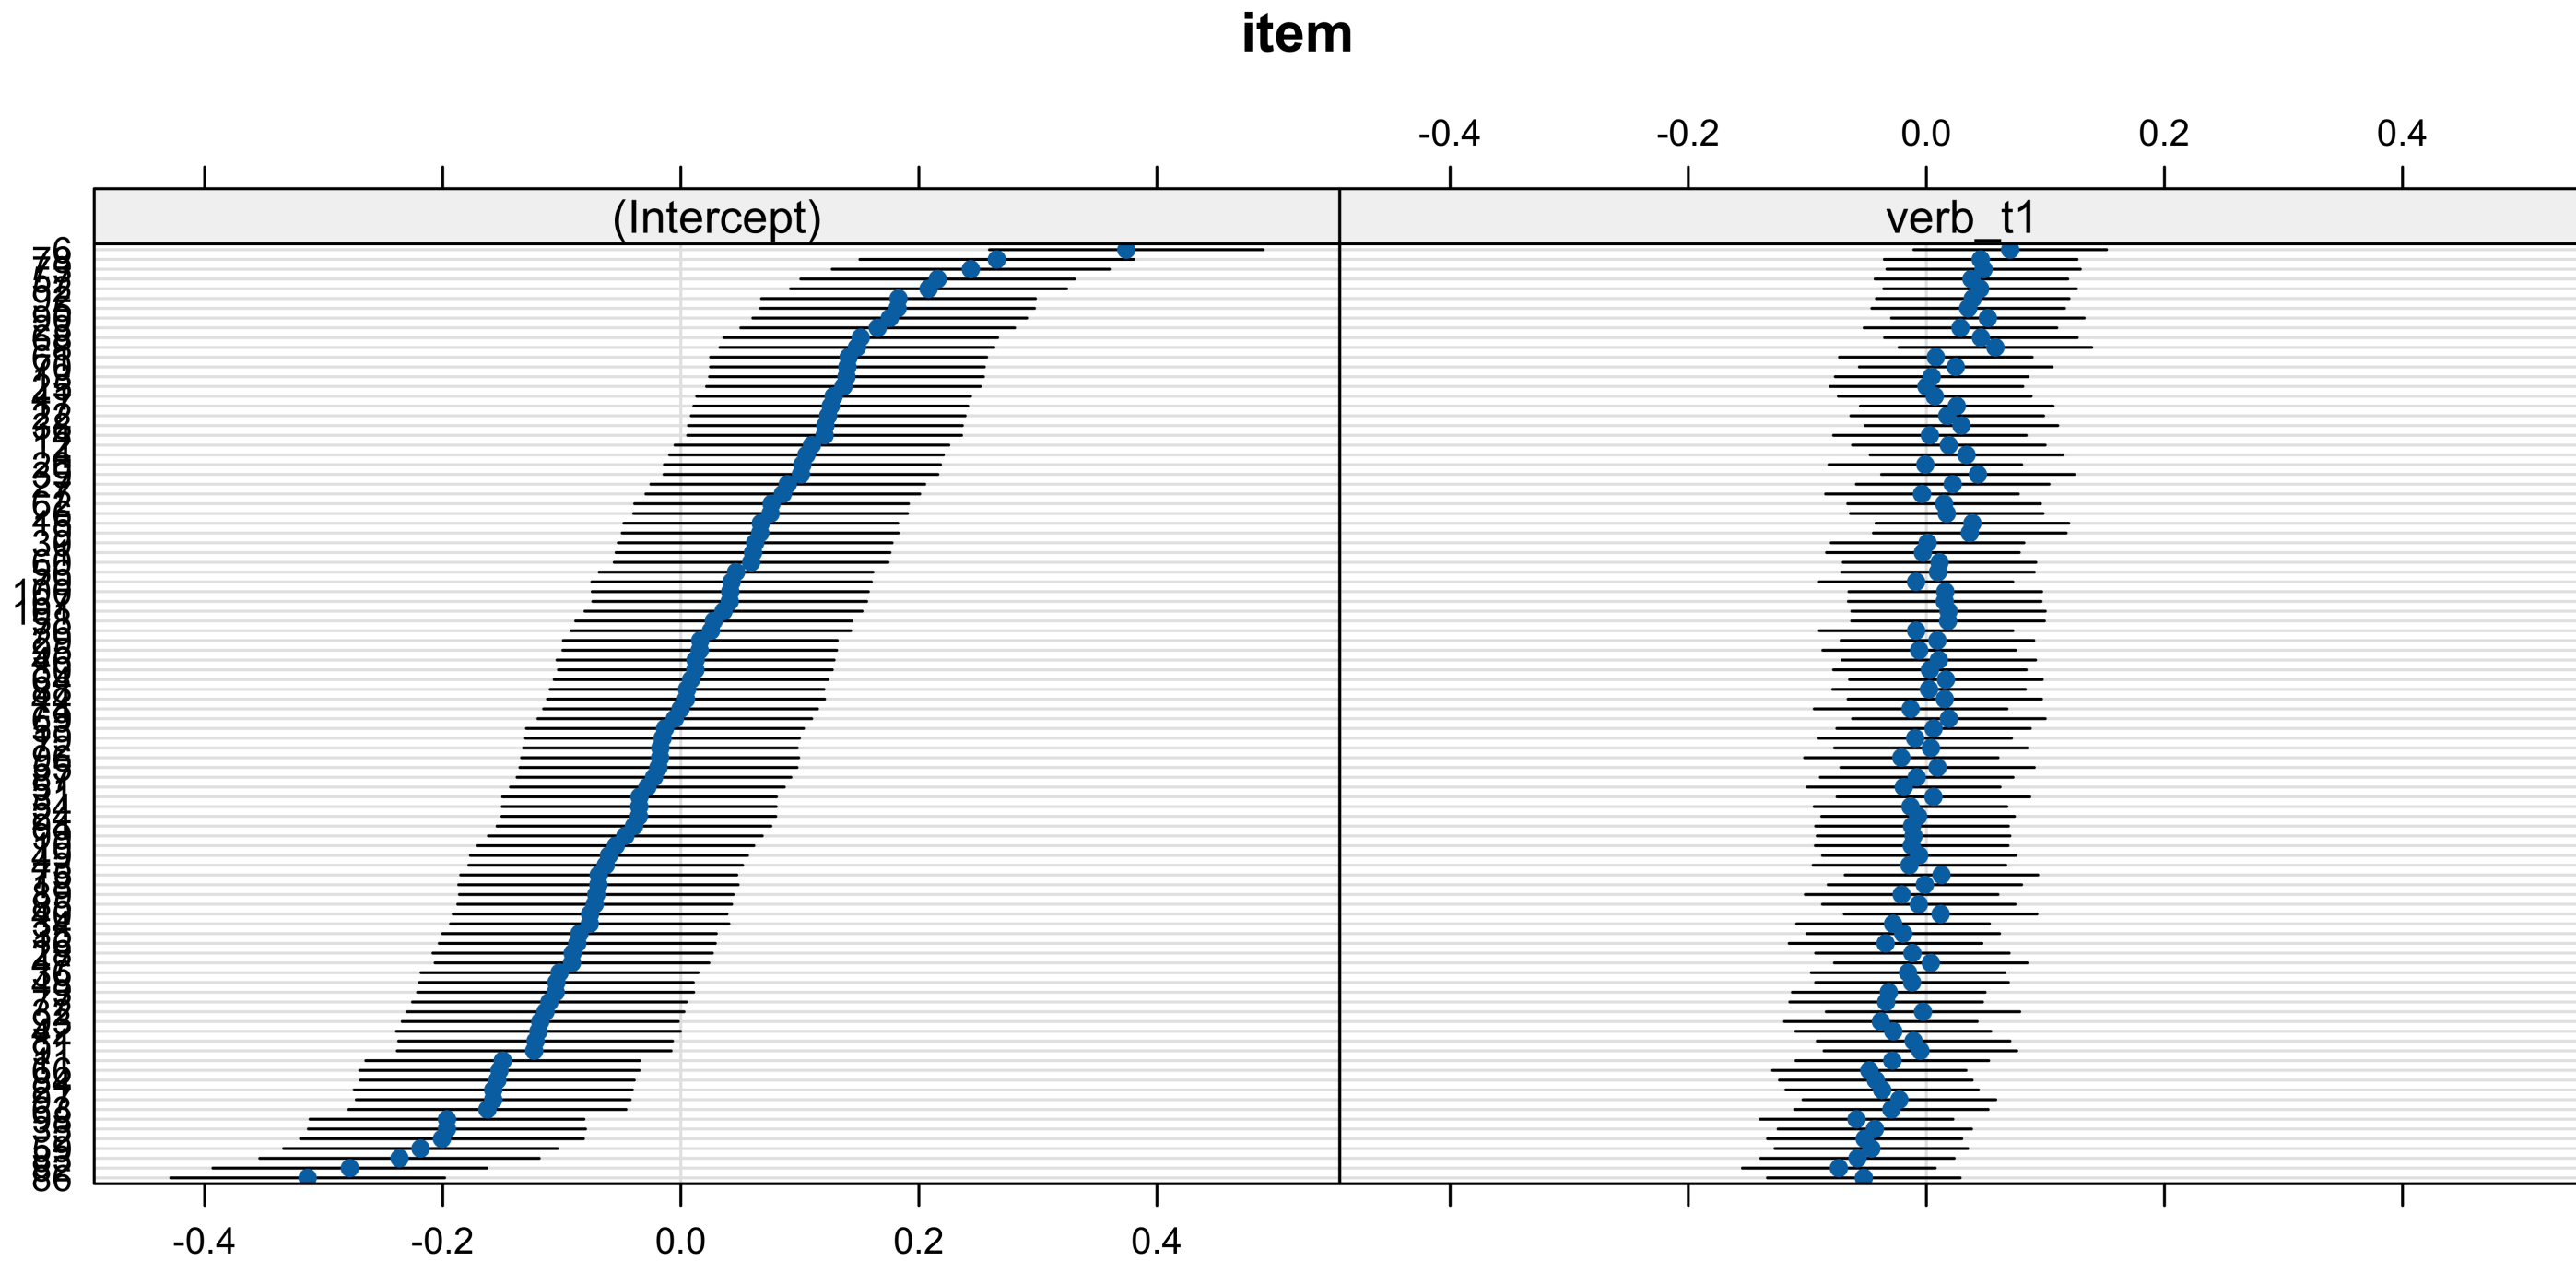
by-item random effects

```
1 lattice::dotplot(ranef(fit_verb_fp))$sj
```



by-participant random effects

```
1 lattice::dotplot(ranef(fit_verb_fp))$item
```



summary()

```
1 summary(fit_verb_fp)
```

```
Linear mixed model fit by REML. t-tests use Satterthwaite's method [
lmerModLmerTest]
Formula: log(fp) ~ verb_t * gramm + (1 | sj) + (1 + verb_t | item)
  Data: df_biondo
Control: lmerControl(optimizer = "bobyqa", optCtrl = list(maxfun = 200000))
Subset: roi == 4
```

```
REML criterion at convergence: 4216.2
```

```
Scaled residuals:
```

Min	1Q	Median	3Q	Max
-4.1758	-0.6096	-0.0227	0.6060	4.0568

```
Random effects:
```

Groups	Name	Variance	Std.Dev.	Corr
Item	(Intercept)	0.010400	0.10200	
	verb_t	0.000000	0.00000	

- IMPORTANTLY, only look at the fixed effects after you've got your final model!!!!
 - i.e., run model -> convergence error -> `rePCA()` + `VarCorr()` -> run model -> ... -> converges -> only NOW run `summary(model)`

Comparing to ‘bad’ models

- let's compare our final model to our ‘bad’ models
 - random intercepts-only model (overconfident)
 - maximal model (underconfident)

Random-intercepts only

```
1 fit_verb_fp_intercepts <- lmer(log(fp) ~ verb_t*gramm +  
2                               (1 | sj) +  
3                               (1 | item),  
4                               data = df_biondo,  
5                               subset = roi == 4  
6 )
```

- converges

► Code

coefficient estimates

► Code

Table 1: Coefficient estimates for our parsimonious model, a random-intercepts only model, and a maximal model

term	parsimonious	intercepts	maximal	measure
(Intercept)	5.9564	5.9564	5.9564	estimate
verb_t1	0.0617	0.0619	0.0617	estimate
gramm1	0.0033	0.0032	0.0034	estimate
verb_t1:gramm1	-0.0144	-0.0143	-0.0142	estimate

standard error

► Code

Table 2: Standard error of coefficient estimates for our parsimonious model, a random-intercepts only model, and a maximal model

term	parsimonious	intercepts	maximal	measure
(Intercept)	0.0368	0.0368	0.0367	std.error
verb_t1	0.0140	0.0130	0.0144	std.error
gramm1	0.0130	0.0130	0.0133	std.error
verb_t1:gramm1	0.0260	0.0260	0.0278	std.error

- standard error ($SE = \frac{\sigma}{\sqrt{n}}$) is a measure of uncertainty
 - larger values reflect greater uncertainty
 - because n is in the denominator, SE gets smaller with more observations
- compared to our parsimonious model with by-item varying **verb_t** slopes:
 - smaller SE for our overconfident (intercepts) model
 - larger SE for our underconfident (maximal) model
 - but only for the estimate also included in the random effects

t-values

► Code

Table 3: t-values of each estimates for our parsimonious model, a random-intercepts only model, and a maximal model

term	parsimonious	intercepts	maximal	measure
(Intercept)	162.0213	161.9025	162.1605	statistic
verb_t1	4.4188	4.7517	4.2982	statistic
gramm1	0.2537	0.2466	0.2542	statistic
verb_t1:gramm1	-0.5531	-0.5496	-0.5108	statistic

- t-value ($t = \frac{\bar{x}_1 - \bar{x}_2}{SE}$) is a measure of uncertainty
 - larger values reflect greater effect
 - more n increases t
- again, **verb_t**: $t_{\max} < t_{\text{pars}} < t_{\text{int}}$

degrees of freedom

► Code

Table 4: Degrees of freedom of each estimates for our parsimonious model, a random-intercepts only model, and a maximal model

term	parsimonious	intercepts	maximal	measure
(Intercept)	79.2432	79.2008	79.1789	df
verb_t1	93.4106	3637.1332	71.4326	df
gramm1	3544.4518	3637.1834	180.0819	df
verb_t1:gramm1	3544.7623	3637.1023	91.8570	df

- degrees of freedom: not trivially defined in mixed models; we're using Satterthwaite approximation (default in `lmerTest::lmer()`)
 - larger degrees of freedom corresponds to larger `n`
 - including more random effects reduces our `n` and therefore reduces `df`
- again, `verb_t`: $df_{\text{max}} < df_{\text{pars}} < df_{\text{int}}$
 - and large differences between our maximal model and the other two for other terms

p-values

► Code

Table 5: p-values of coefficient estimates for our parsimonious model, a random-intercepts only model, and a maximal model

term	parsimonious	intercepts	maximal	measure
(Intercept)	0.0000000	0.0000000	0.0000000	p.value
verb_t1	0.0000267	0.0000021	0.0000535	p.value
gramm1	0.7997645	0.8052568	0.7996177	p.value
verb_t1:gramm1	0.5802114	0.5826522	0.6107494	p.value

- p-values: inversely related to t-values (larger t-values = smaller p-values)
- again, **verb_t**: $p_{\text{max}} < p_{\text{pars}} < p_{\text{int}}$
 - this would be important for ‘significance’ if the values were closer to the conventional alpha-levels ($p < .05$, $p < .01$, $p < .001$)
 - but here the different random effects structures don’t qualitatively change (all are $< .001$)
- this is not always the case, however!
 - this is why we do not peek at the fixed effects until we have our final model
 - we don’t want to be influenced (consciously or not) by seeing small p-values in one model but not another

Reporting

- in Data Analysis section, e.g.,

We included Time Reference (past, future), and Verb Match (match, mismatch) as fixed-effect factors in the models used to investigate the processing of past–future violations (Q1), by adopting sum contrast coding (Schad et al., 2020): past and match conditions were coded as $-.5$, while future and mismatch conditions were coded as $.5$. [...] Moreover, we included crossed random intercepts and random slopes for all fixed-effect parameters for subject and item grouping factors (Barr et al., 2013) in all models.

We reduced the complexity of the random effect structure of the maximal model by performing a principal component analysis so as to identify the most parsimonious model properly supported by the data (Bates et al., 2015). [...] all reading time data were log transformed before performing the analyses.

— Biondo et al. (2022), p. 9

Formatted p-values

- we can use the `format_pval()` function defined earlier to produce formatted p-values




```
1 tidy(fit_verb_fp,  
2     effects = "fixed") |>  
3 as_tibble() |>  
4 mutate(p_value = format_pval(p.value)) |>  
5 select(-p.value) |>  
6 kable() |>  
7 kable_styling()
```

Table 6: Table with formatted p-values from `format_pval()`

effect	term	estimate	std.error	statistic	df	p_value
fixed	(Intercept)	5.9563839	0.0367630	162.0213386	79.24318	< .001
fixed	verb_t1	0.0617330	0.0139706	4.4187860	93.41060	< .001
fixed	gramm1	0.0032976	0.0129994	0.2536709	3544.45182	0.800
fixed	verb_t1:gramm1	-0.0143804	0.0259984	-0.5531269	3544.76235	0.580

Learning objectives

Today we...

- applied remedies for nonconvergence 
- reduced our RES with a data-driven approach 
- compared a parsimonious model to maximal and intercept-only models 

Important terms

Term	Definition	Equation/Code
linear mixed (effects) model	NA	NA

References

- Barr, D. J., Levy, R., Scheepers, C., & Tily, H. J. (2013). Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language*, 68(3), 255–278. <https://doi.org/10.1016/j.jml.2012.11.001>
- Bates, D., Kliegl, R., Vasishth, S., & Baayen, H. (2015). Parsimonious Mixed Models. *arXiv Preprint*, 1–27. <https://doi.org/10.48550/arXiv.1506.04967>
- Biondo, N., Soilemezidi, M., & Mancini, S. (2022). Yesterday is history, tomorrow is a mystery: An eye-tracking investigation of the processing of past and future time reference during sentence reading. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 48(7), 1001–1018. <https://doi.org/10.1037/xlm0001053>
- Brauer, M., & Curtin, J. J. (2018). Linear mixed-effects models and the analysis of nonindependent data: A unified framework to analyze categorical and continuous independent variables that vary within-subjects and/or within-items. *Psychological Methods*, 23(3), 389–411. <https://doi.org/10.1037/met0000159>
- Meteyard, L., & Davies, R. A. I. (2020). Best practice guidance for linear mixed-effects models in psychological science. *Journal of Memory and Language*, 112, 104092. <https://doi.org/10.1016/j.jml.2020.104092>
- Sonderegger, M. (2023). *Regression Modeling for Linguistic Data*.
- Winter, B. (2019). Statistics for Linguists: An Introduction Using R. In *Statistics for Linguists: An Introduction Using R*. Routledge. <https://doi.org/10.4324/9781315165547>

